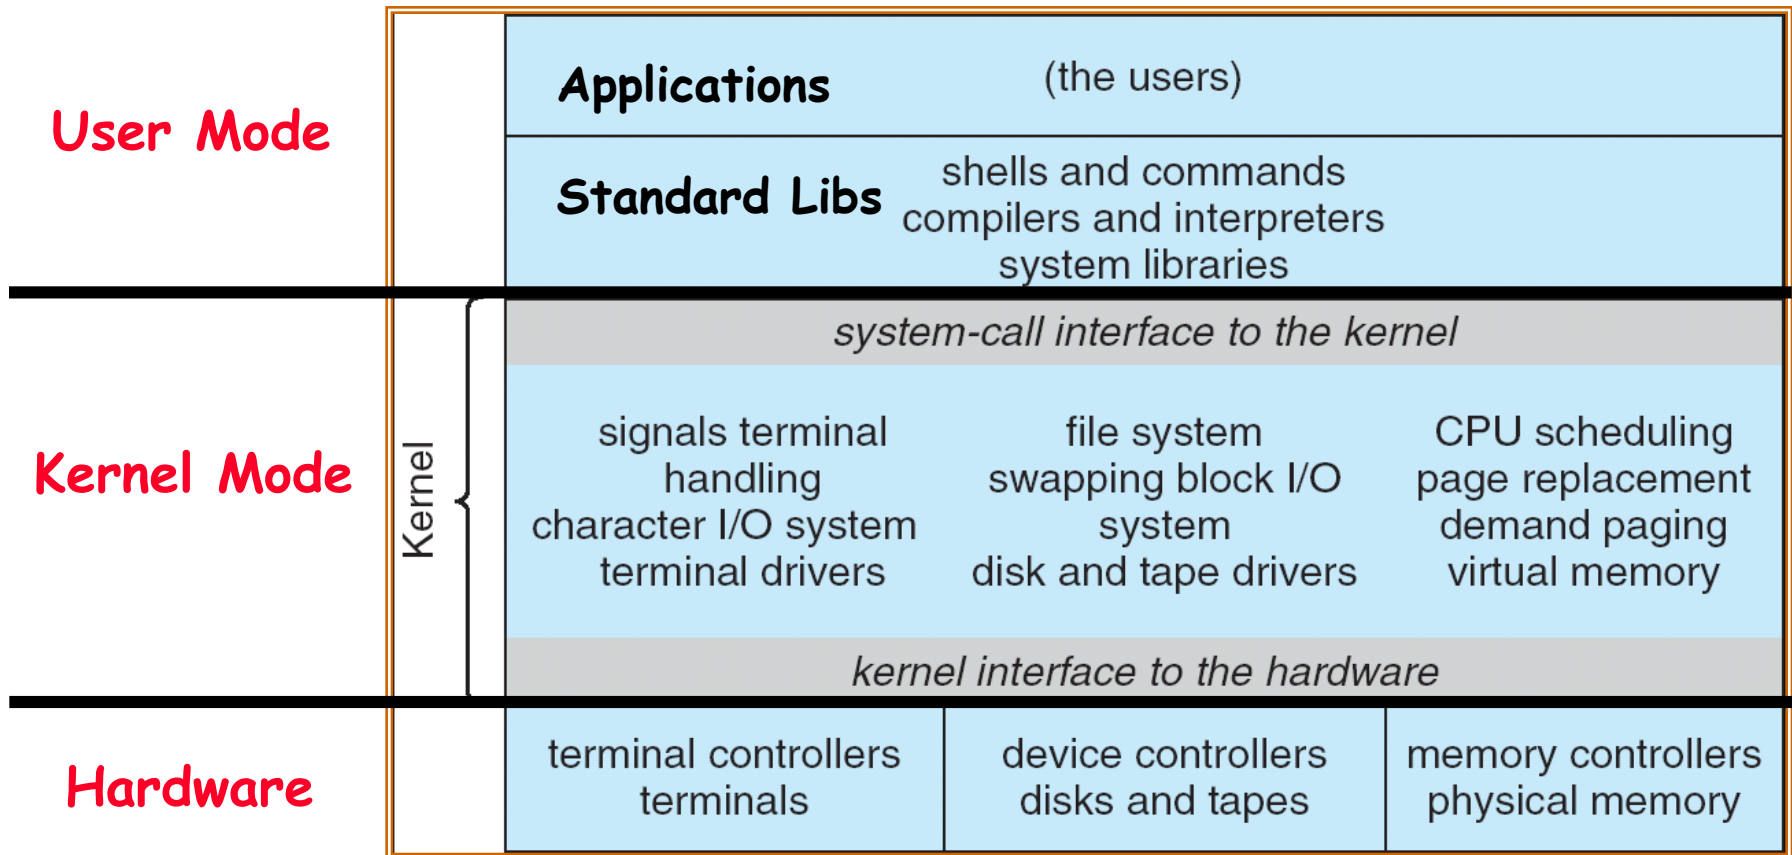


CS3510
Operating Systems
System Calls Interface

Bheemarjuna Reddy Tamma
IIT HYD

UNIX System Structure



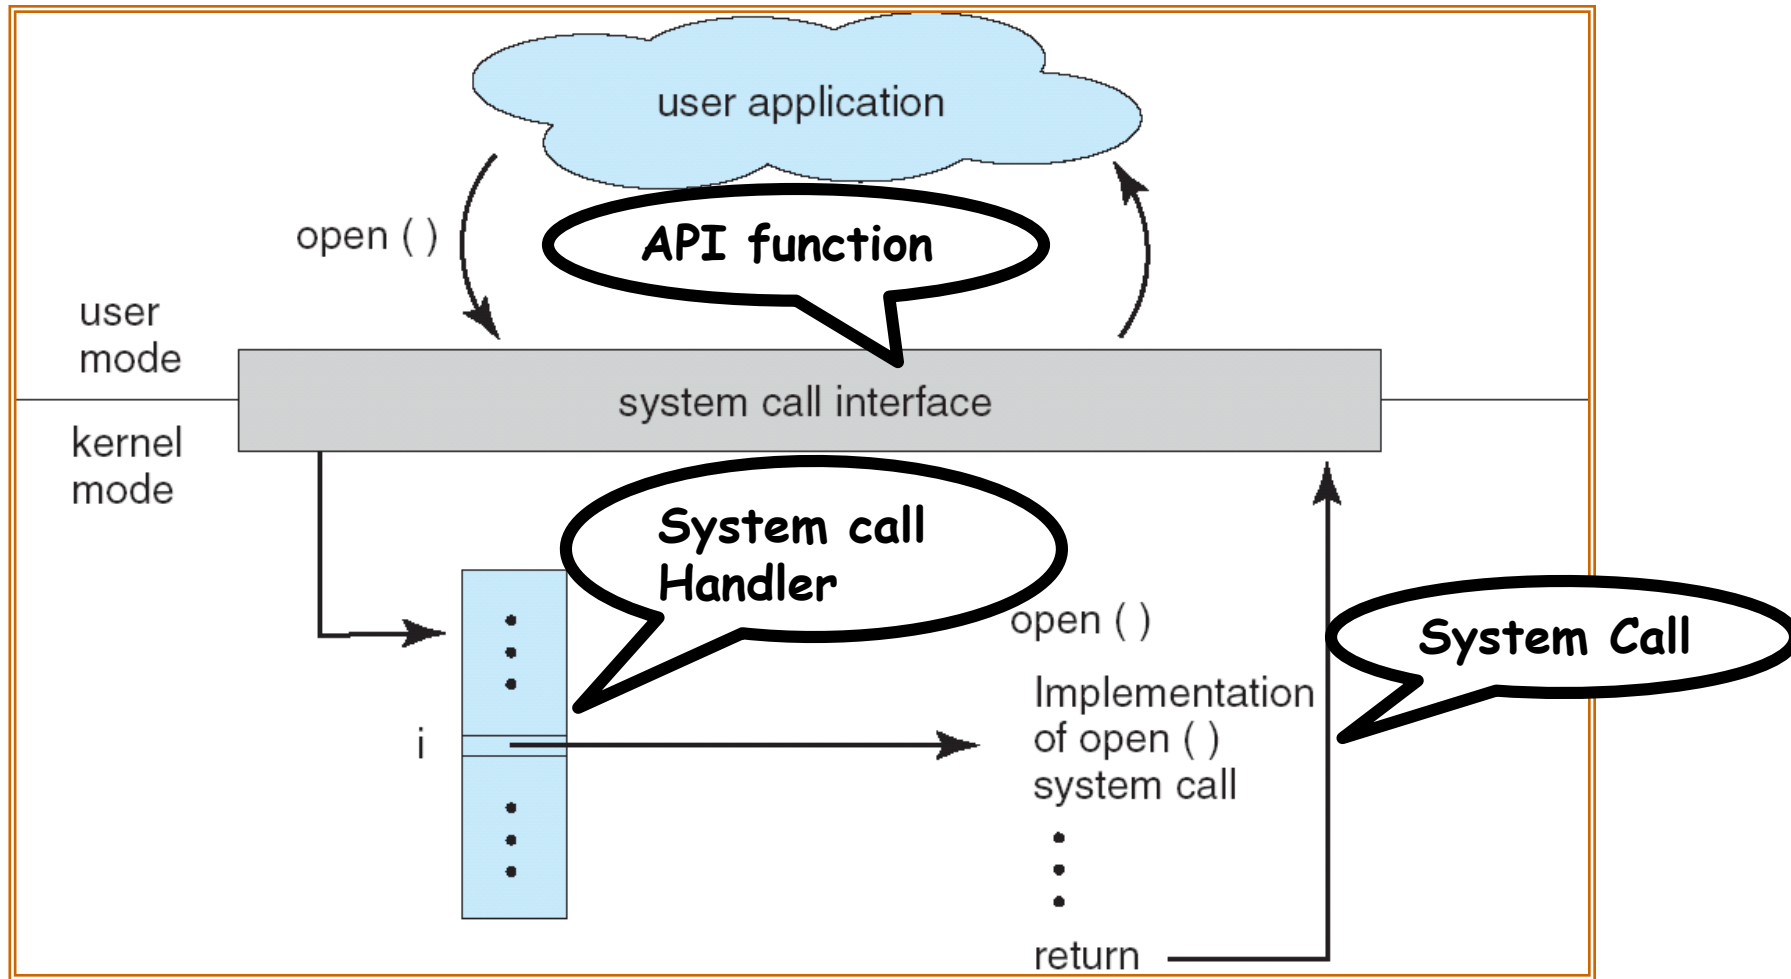
System Calls

- Programming interface to the services provided by the OS to system/app programs
- Typically written in a high-level language (C/C++)
- Mostly accessed by application/system programs using **procedure calls in APIs**
- Programmer/job → procedure in API → System call
 - API is a function definition that specifies how to obtain a given service
 - System call is an explicit request to kernel made via a trap i.e., software interrupt
- Three most common APIs:
 - Win32 API for Windows
 - POSIX API for POSIX-based systems (UNIX, Linux, Mac OS X)
 - Java API for the Java virtual machine (JVM)
- A programmer accesses an API via a library of code (eg., **libc** for C programs in Linux) provided by OS

System Calls

- **POSIX.1-2017** (IEEE 1003.1, ISO/IEC 9945)
 - Very widely used standard based on (and including) C-language
 - POSIX std refers to the API, not actual system calls provided by the kernel
 - Defines both
 - ▶ *API* and
 - ▶ *compulsory system programs/common utilities* together with their functionality and command-line format
 - E.g. `ls -w dir` prints the list of files in a directory in a 'wide' format
 - Complete specification is at <http://www.opengroup.org/onlinepubs/9699919799/nframe.html>
- Strong correlation b/w a procedure/function in API and its associated system call within kernel
 - Typically One-to-one (e.g., read, open, write, exit)
 - many-to-one (e.g., exec, brk) and one-to-many

System Calls



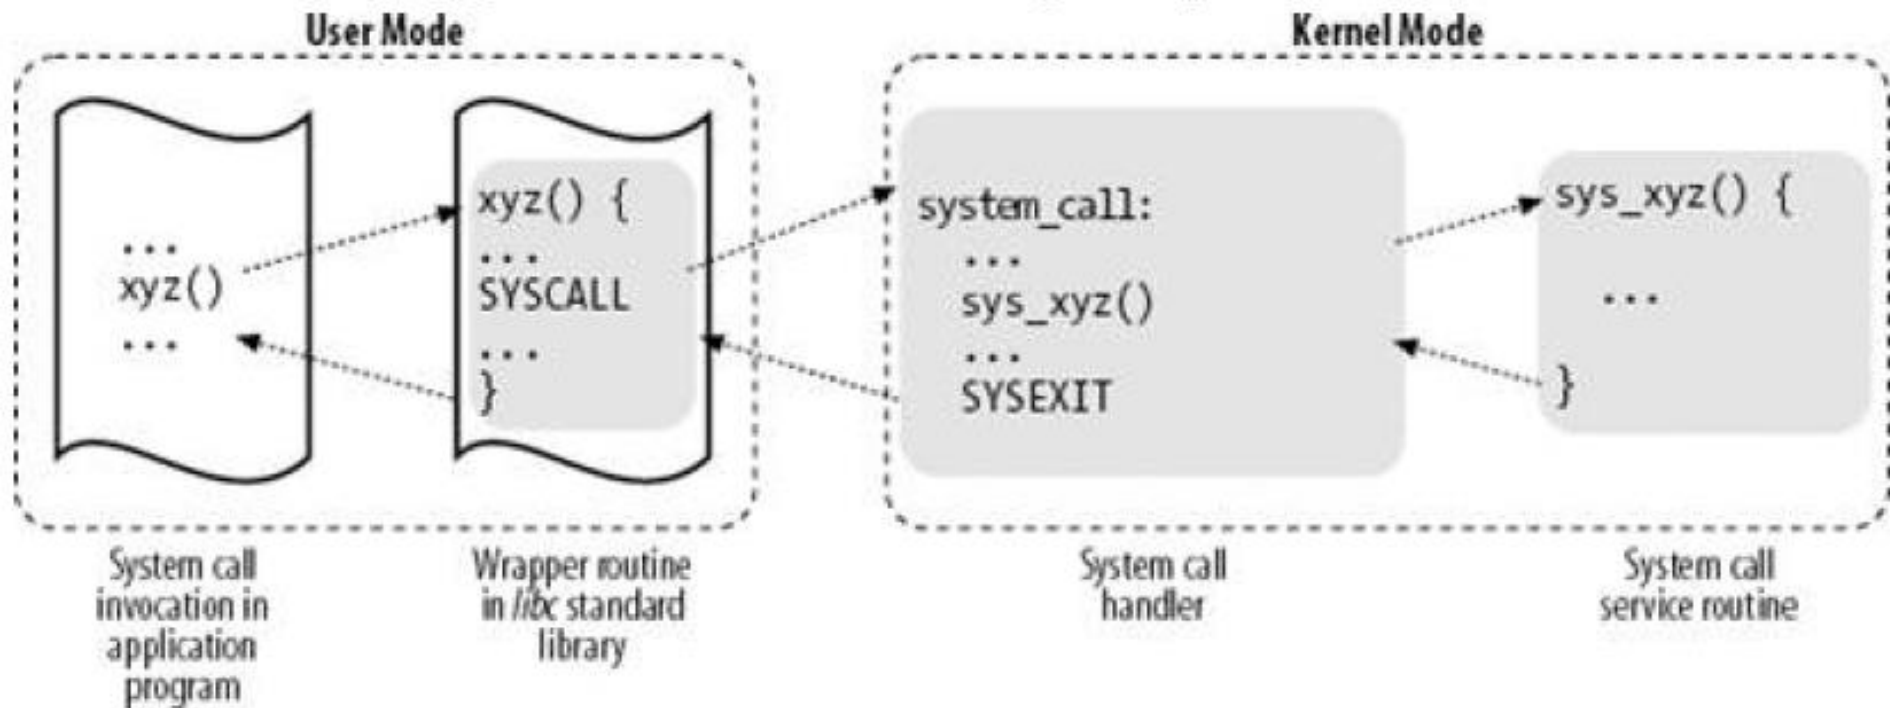
Instructions used to transition to kernel mode in diff archs

- i386 (**int 0x80**), `eax` register is used to indicate syscall number
- x86_64 (**syscall**), `rax` register is used similarly

API-System Call Implementation

- The interface to the services provided by the OS has two parts:
 1. Higher language interface - a part of system library
 - Executes in user mode
 - Implemented to accept standard procedure calls
 - Traps to the Part 2
 2. Kernel part
 - Executes in kernel mode
 - Implements the required system service
 - May cause blocking the caller (forcing it to wait)
 - After completion returns back to Part 1 (may report the success or failure of the call)

System Calls



- Why use APIs rather than system calls directly?
 - Program portability
 - Easier to use

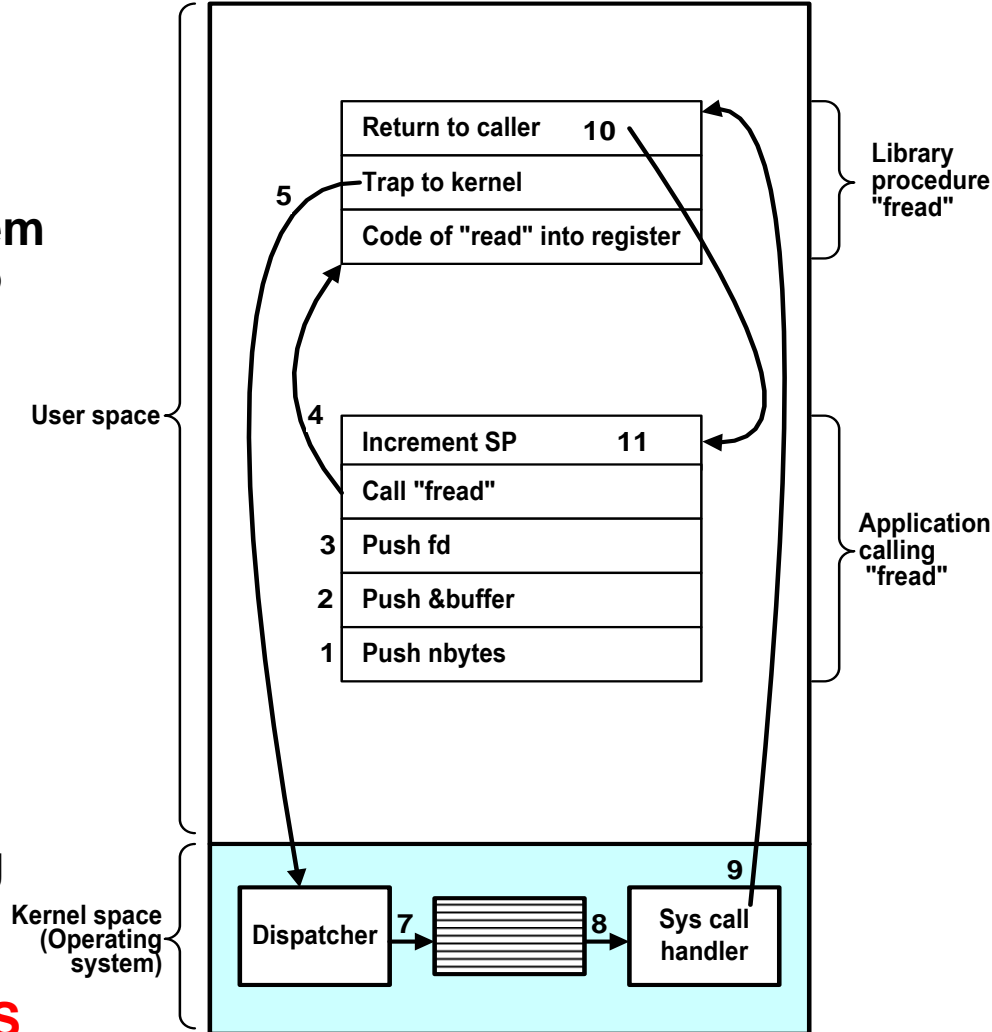
System Call Interface: Implementation

- An application program wants to make use of a System Call:

- A system library routine is called first
- It transforms the call to the system standard (*native API*) and traps to the kernel
- Control is taken by the kernel running in the kernel mode
- According to the service “code”, the *Call dispatcher* invokes the responsible part of the Kernel
- Depending on the nature of the required service, the kernel may block the calling process
- After the call is finished, the calling process execution resumes obtaining the result (success/failure) as if an ordinary function was called

- Three ways to pass parameters to OS

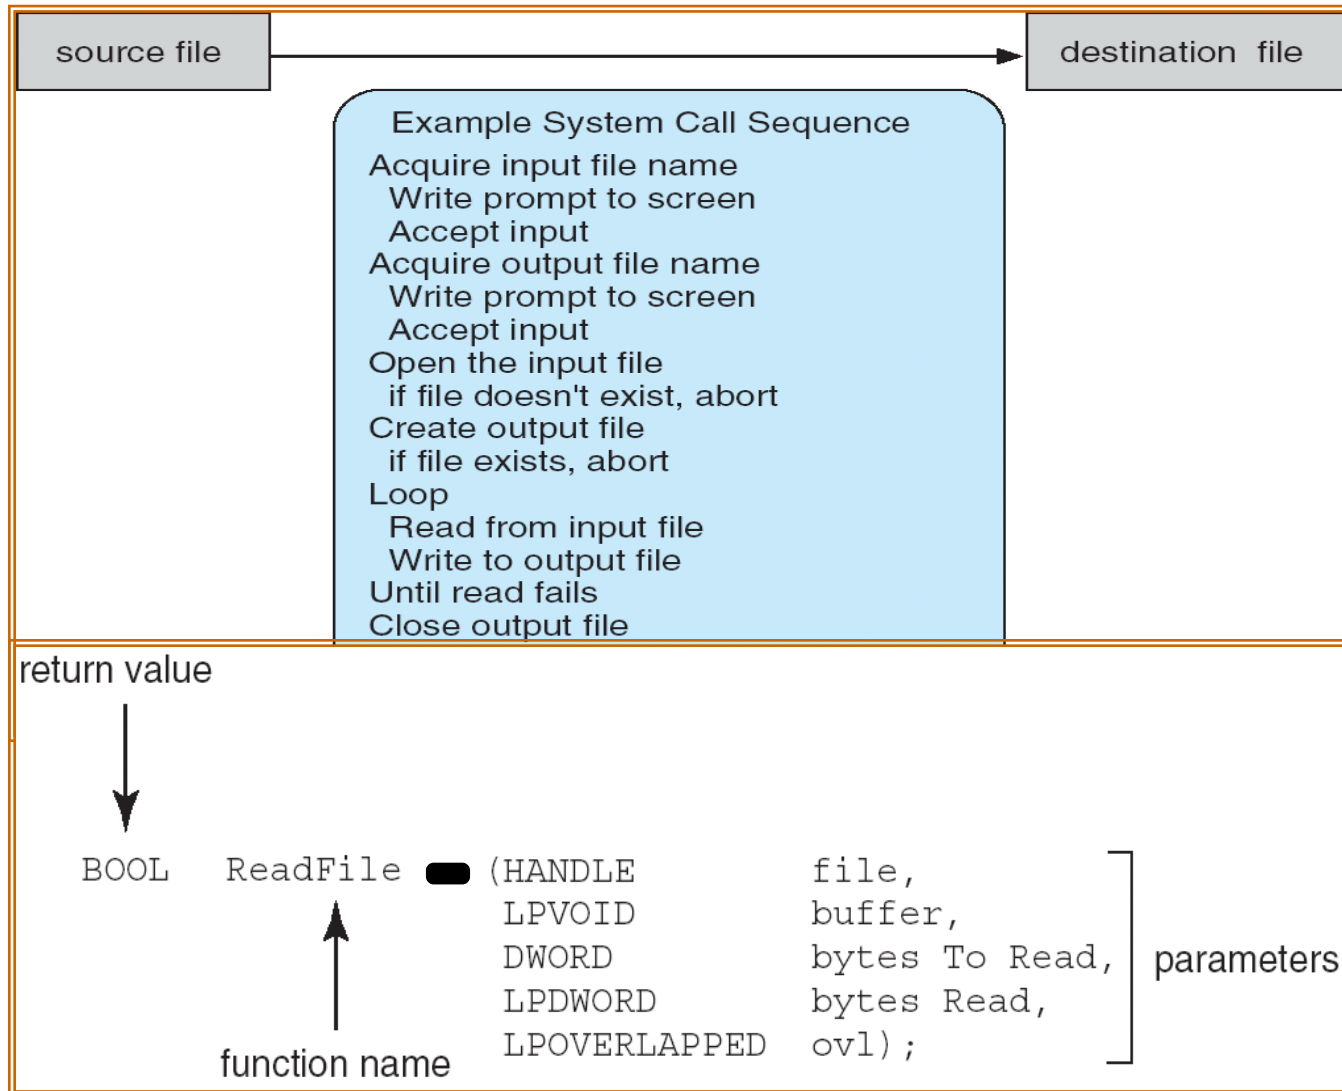
- Registers
- Stack
- Memory block



11 steps to execute the service
fread (fd, buffer, nbytes)

System Calls

- System call sequence to copy contents of one file to another



glibc: GNU C Library

- Any Unix-like OS needs a C library
- C lang has no built-in facilities for doing I/O, memory management, string manipulation, etc
- A std C library (ISO C std) provides these facilities
- The GNU C Library (glibc) implements all of the functions specified in
 - ISO C library (malloc, printf, fopen, exit, etc)
 - POSIX.1 (system calls)
 - And extensions specific to GNU systems
- glibc is used as ***the*** C library in GNU systems and most systems with Linux kernel
 - Current version 2.32 ([link](#))

glibc: GNU C Library

- glibc has procedures (**wrapper functions**) which in turn call system calls
 - getpid(), getppid(), chmod() are defined in glibc
 - glibc provides **syscall** which helps you to call system calls explicitly (directly) from user/app program
 - syscall is also a library function!, but very simple one
 - long syscall (long sysno, ...)
 - sysno is system call number, refer <sys/syscall.h> for Macros
 - Return val is the return value of syscall pointed to by sysno
 - » -1 when system call is failed
 - Employing syscall() is useful when invoking a system call that has no wrapper function defined in the C library
-
- <http://man7.org/linux/man-pages/man2/syscall.2.html>
 - <http://man7.org/linux/man-pages/man2/syscalls.2.html>
 - <http://man7.org/linux/man-pages/man7/vdso.7.html>
 - <http://man7.org/linux/man-pages/man7/libc.7.html>

Example 1

```
#define _GNU_SOURCE
#include <unistd.h> //wrapper for syscalls
#include <sys/syscall.h> // loc: /usr/src/include/i386-linux-gnu/bits/syscall.h, defines syscall numbers/Macros
#include <sys/types.h>
#include <stdio.h>

int main(int argc, char *argv[]) {
    pid_t tid;

    tid = syscall(SYS_gettid); //SYS_gettid does not have glibc wrapper function, so calling syscall directly using "syscall" func; refer man
    syscall, man gettid

    printf("TID=%d\n", tid);

    tid = getpid(); //getpid is wrapper function given in glibc

    printf("PID=%d\n", tid);

    tid = getppid(); //getppid is wrapper in glibc

    printf("PPID=%d\n", tid);

    tid = syscall(__NR_getpid); //calling SYSCALL directly

    printf("PID=%d\n", tid);

    tid = syscall(SYS_getpid); //calling SYSCALL directly

    printf("PID=%d\n", tid);

    tid = syscall(__NR_getppid); //calling SYSCALL directly

    printf("PPID=%d\n", tid);

    return 0; }
```

Example 2: (kind of) direct system call

```
#include <unistd.h>
```

```
#include <sys/syscall.h>
```

```
#include <errno.h>
```

```
...
```

```
int rc;
```

```
rc = syscall(SYS_chmod, "/etc/passwd", 0444);
```

```
if (rc == -1)
```

```
fprintf(stderr, "chmod failed, errno = %d\n", errno);
```

```
...
```

Example 2': glibc wrapper call

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <errno.h>
```

```
...
```

```
int rc;
```

```
rc = chmod("/etc/passwd", 0444);
```

```
if (rc == -1)
```

```
fprintf(stderr, "chmod failed, errno = %d\n", errno);
```

```
...
```

Some API Calls For Process Management

Process management

Call	Description
<code>pid = fork()</code>	Create a child process identical to the parent
<code>pid = waitpid(pid, &statloc, options)</code>	Wait for a child to terminate
<code>s = execve(name, argv, environp)</code>	Replace a process' core image
<code>exit(status)</code>	Terminate process execution and return status

Some API Calls For File Management

File management

Call	Description
<code>fd = open(file, how, ...)</code>	Open a file for reading, writing or both
<code>s = close(fd)</code>	Close an open file
<code>n = read(fd, buffer, nbytes)</code>	Read data from a file into a buffer
<code>n = write(fd, buffer, nbytes)</code>	Write data from a buffer into a file
<code>position = lseek(fd, offset, whence)</code>	Move the file pointer
<code>s = stat(name, &buf)</code>	Get a file's status information

Some API Calls For Directory Management

Directory and file system management

Call	Description
<code>s = mkdir(name, mode)</code>	Create a new directory
<code>s = rmdir(name)</code>	Remove an empty directory
<code>s = link(name1, name2)</code>	Create a new entry, name2, pointing to name1
<code>s = unlink(name)</code>	Remove a directory entry
<code>s = mount(special, name, flag)</code>	Mount a file system
<code>s = umount(special)</code>	Unmount a file system

Some API Calls For Other Tasks

Miscellaneous

Call	Description
<code>s = chdir(dirname)</code>	Change the working directory
<code>s = chmod(name, mode)</code>	Change a file's protection bits
<code>s = kill(pid, signal)</code>	Send a signal to a process
<code>seconds = time(&seconds)</code>	Get the elapsed time since Jan. 1, 1970

POSIX and Win32 Calls Comparison

■ Only some important calls are shown

POSIX	Win32	Description
fork	CreateProcess	Create a new process
wait	WaitForSingleObject	The parent process may wait for the child to finish
execve	--	CreateProcess = fork + execve
exit	ExitProcess	Terminate process
open	CreateFile	Create a new file or open an existing file
close	CloseHandle	Close a file
read	ReadFile	Read data from an open file
write	WriteFile	Write data into an open file
lseek	SetFilePointer	Move read/write offset in a file (file pointer)
stat	GetFileAttributesExt	Get information on a file
mkdir	CreateDirectory	Create a file directory
rmdir	RemoveDirectory	Remove a file directory
link	--	Win32 does not support “links” in the file system
unlink	DeleteFile	Delete an existing file
chdir	SetCurrentDirectory	Change working directory
chmod	SeFileSecurity	Change file mode bits (rwx)

Hands-on on System calls

- `helloWorld.c`
 - Compile it and save as `helloWorld.o`
- `ltrace ./helloWorld.o` (options `-c`, `-S`, `-t`, `-T`)
- `strace ./helloWorld.o` (options `-c`, `-C`, `-t`, `-w`)
- `time ./helloWorld.o`
- `/usr/bin/time -v ./helloWorld.o`
- `getpid`: procedure call and system call
- `syscall`: to directly invoke system calls from user-space programs
- virtual system calls (e.g. `vDSO gettimeofday()`)

Reducing System Call Overhead

- Problem: User-kernel mode distinction poses performance barrier
 - » Crossing this hardware barrier is costly
 - » System calls take 10x-1000x more time than a regular procedure call
- Solution: Perform some system functionality in user mode itself
 - » by caching results (getpid, gettimeofday)
 - » buffering I/O operations to minimize no. of system calls made (read/write vs. fread/fwrite in API)
 - » *Libraries (DLLs)* can reduce number of system calls
 - E.g., "vDSO" (virtual dynamic shared object) is a small library of read-only type of system calls that the kernel automatically maps into the address space of all user-space applications
 - vDSO on x84_64 offers virtual system call `__vdso_gettimeofday()` for the system call `gettimeofday()`

<https://lwn.net/Articles/771441/> & <https://lwn.net/Articles/615809/>

<http://arkanis.de/weblog/2017-01-05-measurements-of-system-call-performance-and-overhead>

Example

- A stripped down shell:

```
while (TRUE) {                                /* repeat forever */
    type_prompt( );                           /* display prompt */
    read_command (command, parameters)        /* input from terminal */

    if (fork() != 0) {                        /* fork off child process */
        /* Parent code */
        waitpid( -1, &status, 0);            /* wait for child to exit */
    } else {
        /* Child code */
        execve (command, parameters, 0);     /* execute command */
    }
}
```

Reading and Viewing Assignments

- Appendix A from Understanding Linux Kernel by Bovet et al
- <http://www.gnu.org/software/libc/>
- <http://www.gnu.org/software/libc/documentation.html>
- Man syscall, syscalls, intro (man -a intro), libc, etc
- <https://www.kernel.org/doc/man-pages/>
- https://en.wikipedia.org/wiki/Unified_Extensible_Firmware_Interface

*** Professor Messer's Linux+ Training:**

<http://www.youtube.com/playlist?list=PLCDA423AB5CEC8FDB>

<http://www.youtube.com/watch?v=6eTi2qu4Fb0&feature=c4-overview&list=UUkefXKtInZ9PLsoGRtml2FQ>