
CS3510

Threads

Bheemarjuna Reddy Tamma
IIT HYD

Outline

- Case for Threads
- Threads vs Processes
- Thread details
 - Pthread Library

Case for Parallelism

```
main()  
  read_data();  
  for(all data i  $\leftarrow$  1 to N)  
    compute(i);  
    write_data(i);  
  endfor
```

Blocking Write

```
main()  
  read_data();  
  for(all data i  $\leftarrow$  1 to N)  
    compute(i);  
    CreateProcess(write_data(i));  
  endfor
```

Does Writing in
new process

Case for Parallelism

Consider the following code fragment:

```
for(k = 0; k < n; k++)  
    a[k] = b[k] * c[k] + d[k] * e[k];
```

Instead:

```
fn(l, m) {  
    for(k = l; k < m; k++)  
        a[k] = b[k] * c[k] + d[k] * e[k];  
}
```

CreateProcess(fn, 0, n/2);

CreateProcess(fn, n/2, n);

Parallelism vs Concurrency

- Both are not same, but typically used interchangeably!
- **Parallelism: Doing multiple tasks simultaneously**
 - E.g., driving while talking!
 - E.g., Running Word App and Media player simultaneously, each one on separate Core/CPU
 - Not possible w/o multiple Cores
- **Concurrency: Making progress for multiple tasks**
 - Single CPU: time-sharing of CPU resource with time slices
 - Multiple CPU: same as parallelism
- So, Parallelism → Concurrency, but **not vice versa**.

Processes Overheads

- A full process includes numerous things:
 - an address space (defining all the code and data segments)
 - OS resources and accounting information
 - a “thread of control”,
 - defines where the process is currently executing
 - That is the PC and other registers

↳ Creating a new process is costly

- all of the structures (e.g., page tables, address space) that must be allocated
- Costly even after copy-on-write optimization

↳ Communicating between processes (IPC) is costly

- most communication goes through the OS to avoid synchronization issues with shared resources

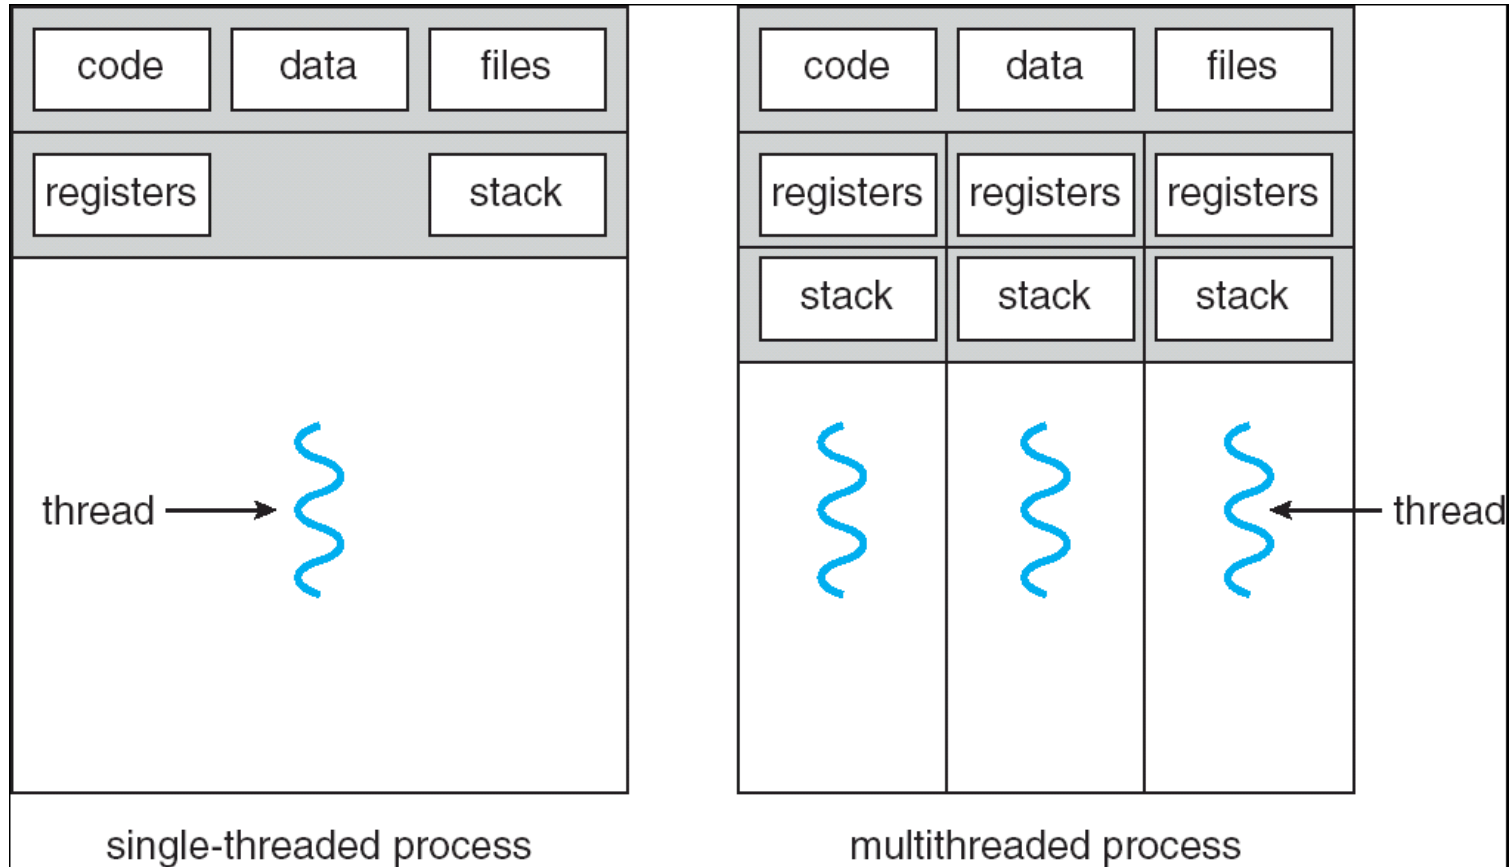
Need “Lightweight” Processes

- What's similar in these processes?
 - They all share the same code, heap and data i.e., most of the address space
 - They all share the same privileges
 - They share almost everything in the process
- What don't they share?
 - Each has its own PC, register set, stack, and stack pointer
- Idea: why don't we separate the idea of process (address space, accounting, etc.) from that of the minimal “thread of control” (PC, SP, registers)?

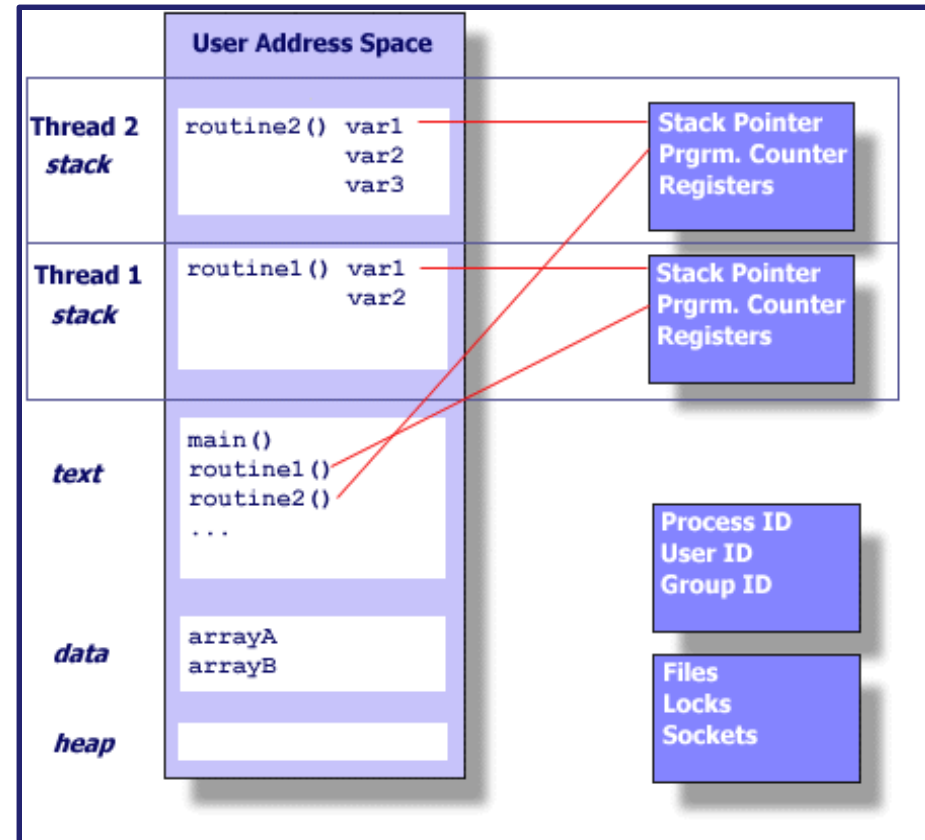
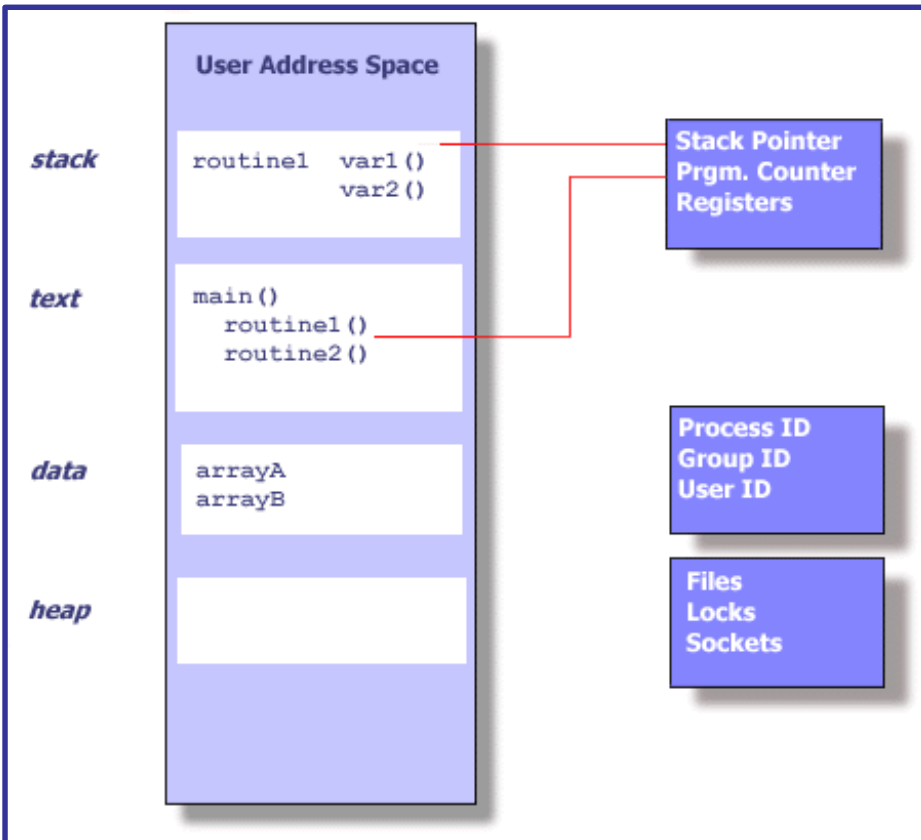
Threads and Processes

- Most operating systems therefore support two entities:
 - the process,
 - which defines the address space and general process attributes
 - the thread,
 - which defines a sequential execution stream within a process
- A thread is bound to a single process.
 - For each process, however, there may be many threads.
- Threads are the unit of scheduling
- Processes are *containers* in which threads execute

Multithreaded Processes



Single vs Multithreaded Processes

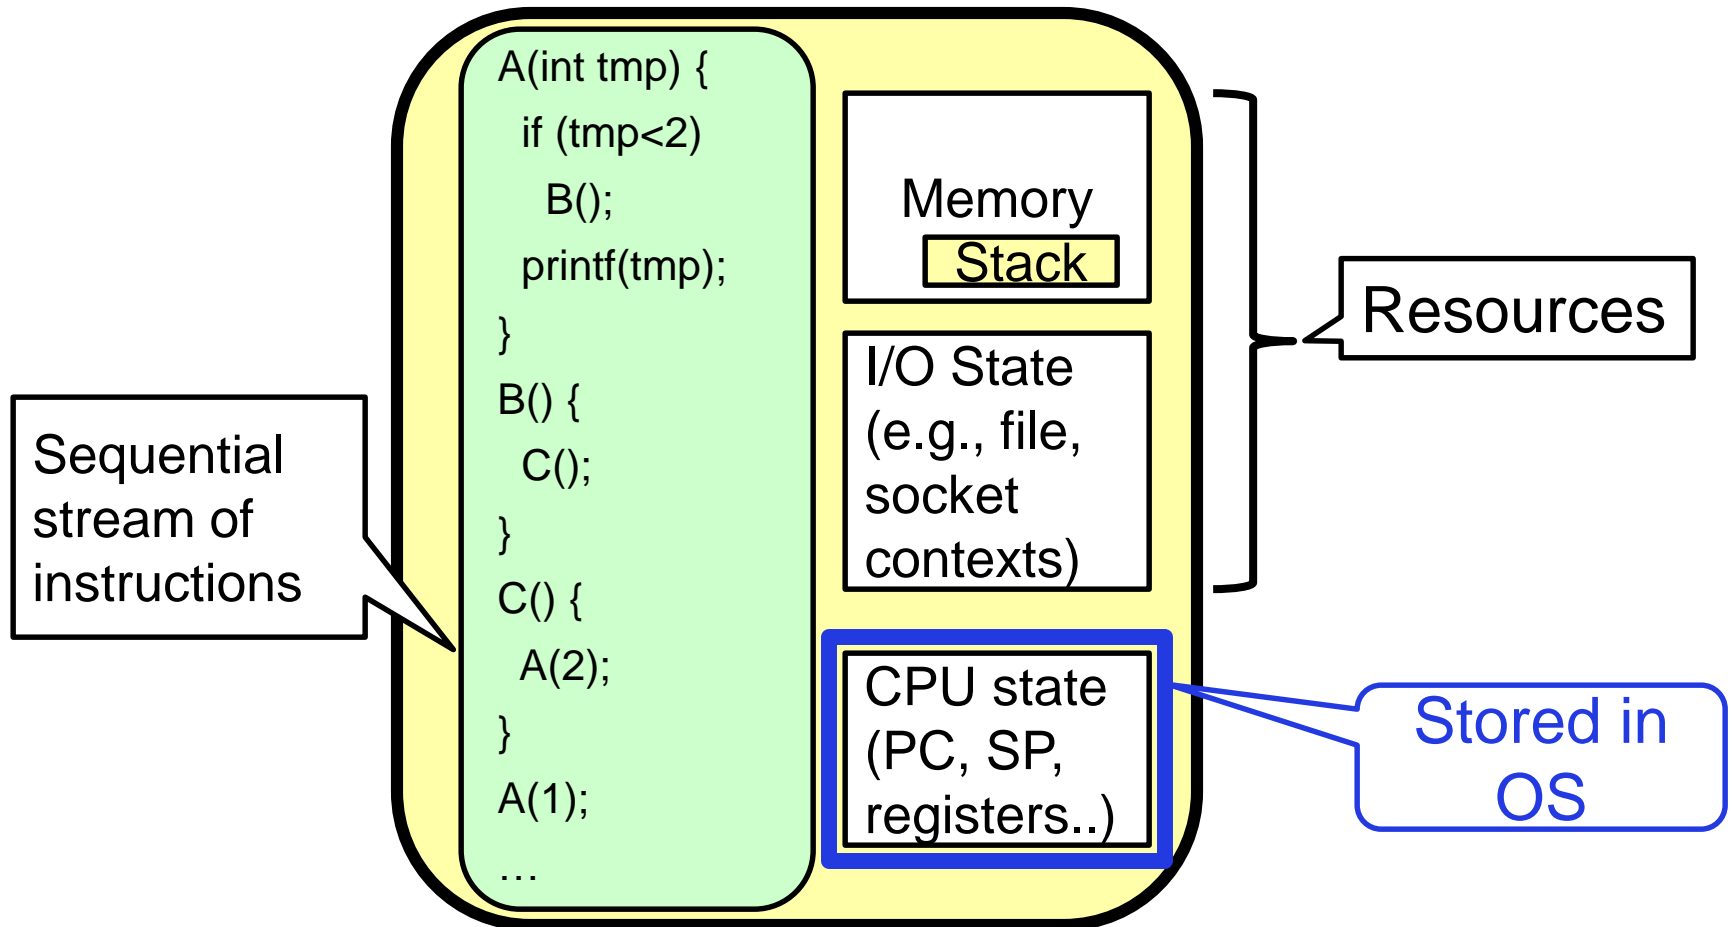


Threads vs. Processes

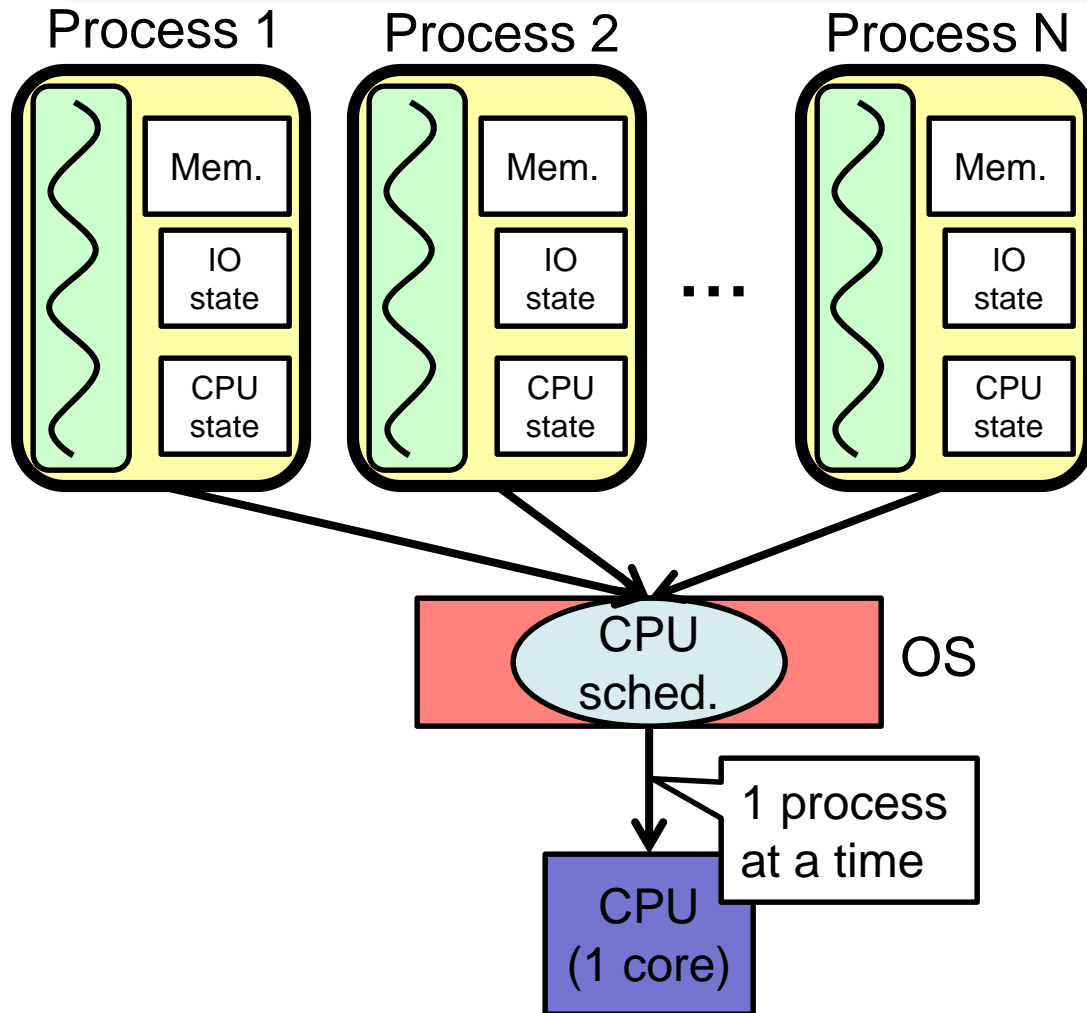
- A thread has no separate code/data segment or heap
- A thread cannot live on its own, it must live within a process
- *There can be more than one thread in a process, the first thread calls main & has the process's stack*
- Inexpensive creation
- Inexpensive context switching
- If a thread dies, its stack is reclaimed
- A process has code/data/heap & other segments
- There must be at least one thread in a process
- *Threads within a process share code/data/heap, share I/O, but each has its own stack & registers*
- Expensive creation
- Expensive context switching
- If a process dies, its resources are reclaimed & all threads die

Putting it Together: Process

(Unix) Process




Putting it Together: Processes



- **Switch overhead:**
 - CPU state: *low*
 - Memory/IO state: *high*
- **Process creation: *high***
- **Protection**
 - CPU: *yes*
 - Memory/IO: *yes*
- **Sharing overhead: *high***

Putting it Together: Threads



- **Switch overhead:**
 - CPU state: *low*
- **Thread creation:** *low*
- **Protection**
 - CPU: *yes*
 - Memory/IO: *no*
- **Sharing overhead:** *low*

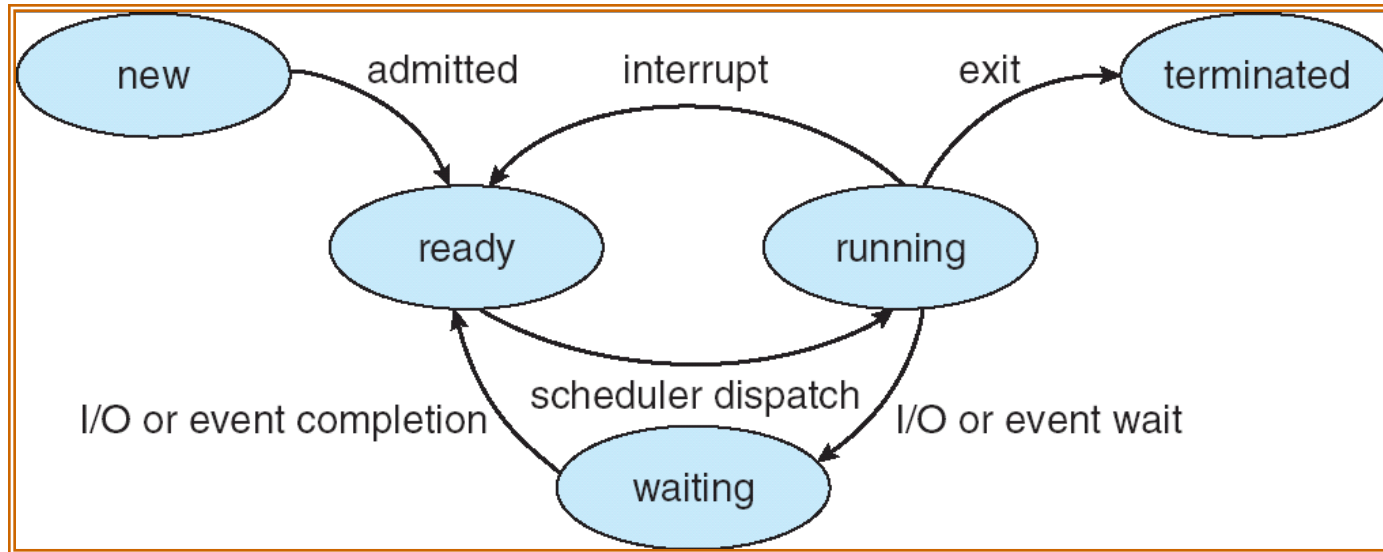
Benefits of multithreaded programs

- **Responsiveness**
 - Interactive apps like Web server
- **Resource sharing**
 - Implicit as threads share the same address space
- **Economy**
 - Creation of multithread process is cheaper
 - Solaris
 - Process creation is 30 times slower than that of thread creation
 - Context switch is about 5 times slower
- **Scalability**
 - Single-threaded process can only run on one CPU
 - Multithreading in multi-core systems increases parallelism

Examples of multithreaded programs

- Embedded systems
 - Elevators, Planes, Medical systems
 - Single Program, but concurrent operations
- Most modern OS kernels
 - Internally concurrent because have to deal with concurrent requests by multiple users
 - Threads for I/O devices, interrupt handling, managing amount of free memory, etc
 - But no protection needed within kernel
- Database Servers
 - Access to shared data by many concurrent users
 - Also background utility processing must be done

Lifecycle of a Thread (or Process)



- As a thread executes, it changes state:
 - **new**: The thread is being created
 - **ready**: The thread is waiting to run
 - **running**: Instructions are being executed
 - **waiting**: Thread waiting for some event to occur
 - **terminated**: The thread has finished execution
- “Active” threads are represented by their TCBs
 - TCBs organized into queues based on their state

Cooperative Threads

A cooperative thread runs until it decides to give up the CPU

```
main()
```

```
{
```

```
    tid t1 = CreateThread(fn, arg);
```

```
    ...
```

```
    Yield(t1);
```

```
}
```

```
fn(int arg)
```

```
{
```

```
    ...
```

```
    Yield(any);
```

```
}
```

Cooperative Threads

- Cooperative threads use **non pre-emptive** scheduling
- Scheduler gets invoked only when Yield is called
- A thread could also yield the CPU when it blocks for I/O
- Advantages:
 - Simple
 - Scientific apps
- Disadvantages:
 - For badly written code

Non-Cooperative Threads

- No explicit control passing among threads
- Rely on the CPU scheduler to decide which thread to run next
- A thread can be pre-empted at any point
- Often called **pre-emptive** threads
- Most modern thread packages use this approach
 - Pthreads API
 - Win32 threads API
 - Java API

Classification of OS

# threads Per AS: # of addr spaces:	One	Many
One	MS/DOS, early Macintosh	Traditional UNIX
Many	Embedded systems (Geoworks, VxWorks, JavaOS, etc) JavaOS, Pilot(PC)	Mach, OS/2, Linux Windows 10, Solaris, HP-UX, OS X

- Most operating systems have either
 - One or many address spaces
 - One or many threads per address space

Example User Thread Interface

`t = thread_fork(initial context)`

creates a new thread of control

`thread_start(t)`

starts the named thread

`thread_yield()`

voluntarily gives up the processor

`thread_stop()`


Stops/pauses the calling thread, also called `thread_block`

`thread_exit()`

terminates the calling thread, also called `thread_destroy`

Multithreading models

- There are actually 2 level of threads:
- Kernel threads:
 - Supported and managed directly by the kernel.
 - Windows, Mac, Linux
- User threads:
 - Supported above the kernel, and without kernel knowledge by user-level threads library.
 - E.g., **POSIX Pthreads API**



1:1 mapping b/w user and kernel threads in Windows & Linux

Pthreads

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- *Specification, not implementation*
- API specifies behavior of the thread library, implementation is up to developers of the library
- May be provided either as user-level or kernel-level threads
- Common in UNIX operating systems (Linux & Mac OS X)
 - Implemented by glibc as Native POSIX Thread Library (NPTL)

Pthreads API

- `pthread_create` (thread id, attr, start_routine, arguments_start_routine)
- `pthread_exit` (status)
- `pthread_cancel` (thread id)
- `pthread_attr_init` (attr)
- `pthread_attr_destroy` (attr)
- `pthread_join` (thread id, status)
- `pthread_detach` (thread id)
- `pthread_attr_setdetachstate` (attr, detachstate)
- `pthread_attr_getdetachstate` (attr, detachstate)
- `pthread_self` (), `pthread_equal` (TID1, TID2)

Pthreads: Creating Threads

- **main()** program comprises a single, default thread
 - All other threads must be explicitly created using **pthread_create**, anywhere in the program
- By default, new thread is created with certain attributes (configurable)
 - pthread_attr_init and pthread_attr_destroy are used to initialize/destroy thread attribute object
 - Other routines are then used to query/set specific attributes in the thread attribute object
 - Detached or joinable state
 - Scheduling policy, Scheduling parameters
 - Stack address, Stack size, etc

Pthread API: Creation and Termination

```
#include <pthread.h> //Implemented by glibc
#include <stdio.h>
#define NUM_THREADS    5
void *PrintHello(void *threadNo)
{
    long tid;
    tid = (long*)threadNo;
    printf("Hello World! It's me, thread #%ld!\n", tid);
    pthread_exit(void *threadNo); //can pass status to other threads
}
```

→ Use gcc & g++ with -pthread flag

→ getrlimit: get user limits of system resources like memory, no of processes, timeslice, etc

→ \$ulimit: get and set user limits of system resources

→ \$cat /proc/[PID]/limits

→ \$sudo cat /proc/[PID]/sched

Pthread API: Creation and Termination

```
int main ( )
{
    pthread_t threads[NUM_THREADS];
    int rc;
    long t; void * status;
    for(t=0; t<NUM_THREADS; t++){
        printf("In main: creating thread %ld\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)&t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n",
rc);
            exit(-1);
        }
    }
    /* main() should wait for thread(s) to finish */
    for(t=0; t<NUM_THREADS; t++)
        pthread_join(threads[t],&status); //collect status
    pthread_exit(NULL);
}
```

Pthread API: Thread Argument Passing (safeway)

```
long *taskids[NUM_THREADS];

for(t=0; t<NUM_THREADS; t++)
{
    taskids[t] = (long *) malloc(sizeof(long));
    *taskids[t] = t;
    printf("Creating thread %ld\n", t);
    rc = pthread_create(&threads[t], NULL, PrintHello, (void *)
taskids[t]);
    ...
}
```

Full Source Code:

https://computing.llnl.gov/tutorials/pthreads/samples/hello_arg1.c

Pthread API: Thread Argument Passing (Unsafe)

```
int rc;
```

```
long t;
```

```
for(t=0; t<NUM_THREADS; t++)
```

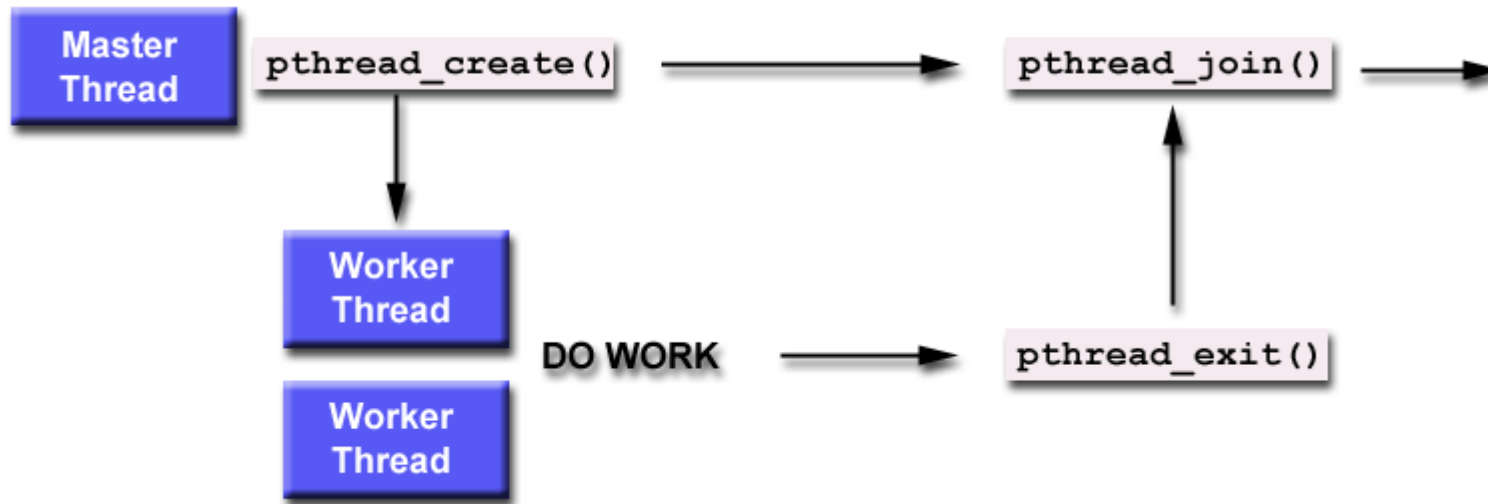
```
{  
    printf("Creating thread %ld\n", t);  
    rc = pthread_create(&threads[t], NULL, PrintHello, (void *) &t);  
    ...  
}
```

- Param **t** is changed by the main thread as it creates new threads

Full Src Code:

https://computing.llnl.gov/tutorials/pthreads/samples/hello_arg3.c

Pthread API: Thread Join/Detach



- "joining" is one of the ways to accomplish synchronization between threads
- Calling `pthread_exit()` at last in `main()` blocks the process till all its threads are done!
- Example:

<https://computing.llnl.gov/tutorials/pthreads/samples/join.c>

Pthread: Issues

After a thread has been created, how do you know

a) when it will be scheduled to run by the OS

b) which processor/core it will run on?

Ans:

1. Depends on underlying thread scheduling algo (FIFO/RO/etc for pthreads) or
2. Implementation specific

Robust programs should not depend on threads running order or core on which a thread runs on

Linux Threads

- Linux does not distinguish between processes and threads
 - Uses term **task** (**struct task_struct**)
 - clone () for creating threads
 - Flags passed as args determine level of sharing b/w parent and child tasks
 - CLONE_FS, CLONE_VM, CLONDE_FILES
 - Sharing → threads
 - No sharing → processes
 - fork() for creating duplicate tasks (processes)

Multithreading Issues

- Semantics of `fork()` and `exec()` system calls
 - Child process duplicates all threads of parent?
 - Two versions of `fork()`!!
 - `exec()` inside a thread will replace the entire process (inc all threads) with prg specified as arg for `exec()`
- Thread cancellation
 - Asynchronous vs. Deferred Cancellation
 - `pthread_cancel` (thread id) supports both, but deferred is recommended as it's safe
- Signal handling
 - Which thread to deliver it to?
 - `kill(pid, signal)`
 - `pthread_kill(tid, signal)`

Thread Hazards

```
int a = 1, b = 2, w = 2;
```

```
main( ) {  
    CreateThread(fn, 4);  
    CreateThread(fn, 4);  
    while(w) ;  
}  
  
fn( ) {  
    int v = a + b;  
    w--;  
}
```

Concurrency Problems

A statement like `w--` in `C` (or `C++`) is implemented by several machine instructions:

```
ld    r4, #w
add   r4, r4, -1
st    r4, #w
```

Now, imagine the following sequence, what is the value of `w`?

```
ld    r4, #w
_____
_____
_____
add   r4, r4, -1
st    r4, #w
```

```
ld    r4, #w
add   r4, r4, -1
st    r4, #w
```

Summary

- Threads increase concurrency/parallelism
- Threads may cause synchronization issues
 - Need to employ synchronization primitives to avoid thread hazards

Reading Assignment

- Chapter 4 from OSC by Galvin et al
- Chapter 2 from MOS by Tanenbaum et al
- <http://www.yolinux.com/TUTORIALS/LinuxTutorialPosixThreads.html>
- <https://computing.llnl.gov/tutorials/pthreads/>