

Database Management Systems (DBMS)

Lec 22: Query processing and optimization (Contd.)

Ramesh K. Jallu

IIIT Raichur

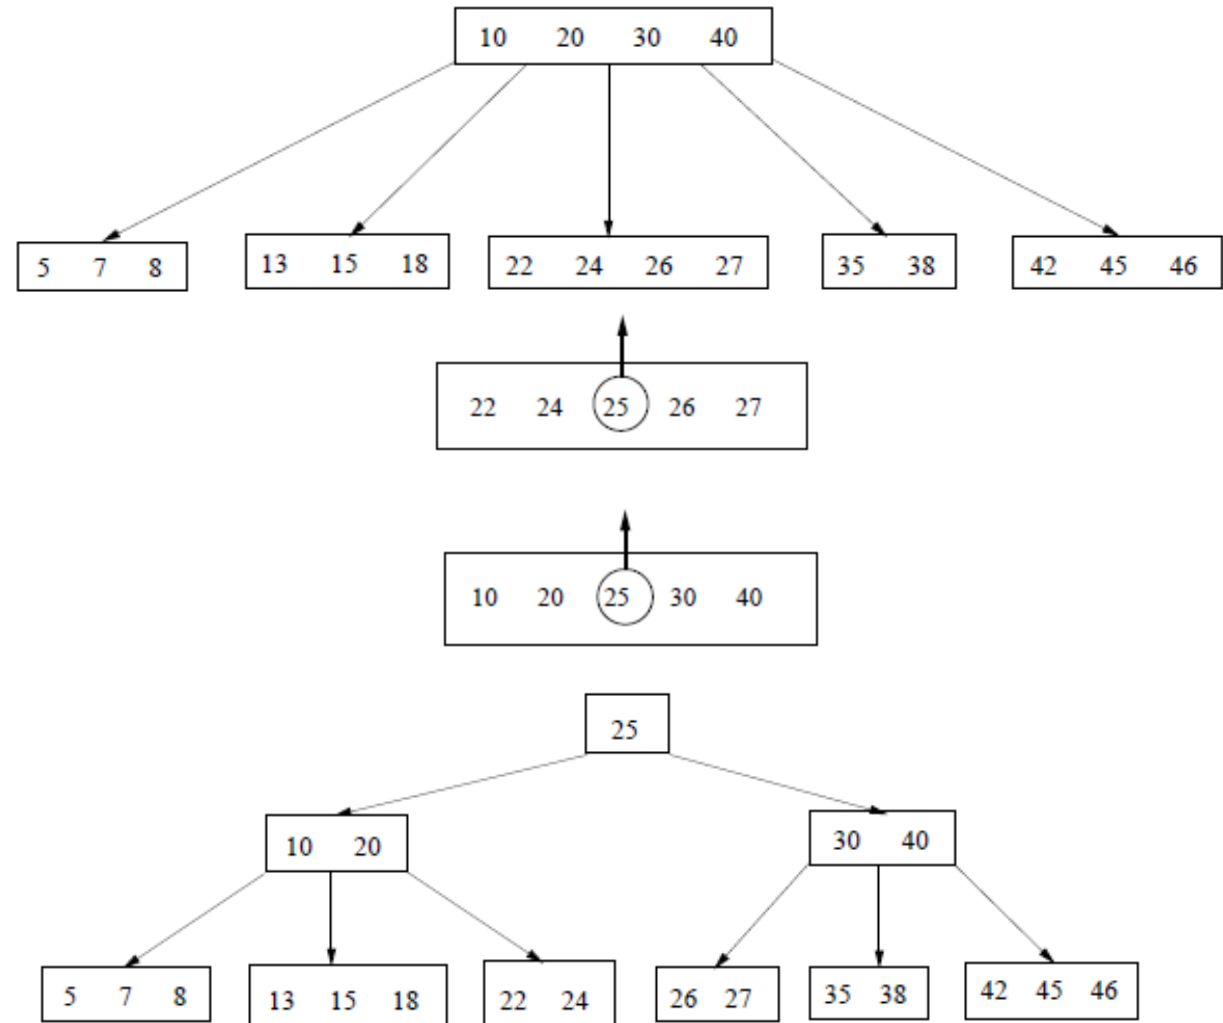


Date: 20/04/21

Recap


- Multilevel indexing
- Search trees
- B-tree

B-tree (Insertion)


* Let $p = 5$; input: 20, 40, 10, 30, 15, 35, 7, 26, 18, 22, 5.




(a) Initial Tree




(b) F deleted: case 1




(c) M deleted: case 2a




(d) G deleted: case 2c




(e) D deleted: case 3b



(e') tree shrinks in height



(f) B deleted: case 3a



B⁺-tree


- In a B-tree, every value of the search key appears at some level in the tree, along with a data pointer
- In a B⁺-tree, data pointers are stored *only at the leaf nodes* of the tree
 - I.e., internal nodes do not store record/block pointers; only leaf nodes do
- The leaf nodes have an entry for *every* value of the search field, along with a data/block pointer to the record if the search field is a *key field*
- For a *nonkey* search field, the pointer points to a block containing pointers to the data file records, creating an extra level of indirection

B⁺-tree (Contd.)

- The leaf nodes of the B⁺-tree are linked to provide ordered access on the search field to the records
- B⁺-tree is used to implement multilevel indexing
 - The leaf nodes are similar to the first (base) level of an index
 - Internal nodes of the B⁺-tree correspond to the other levels
- Some search field values from the leaf nodes are *repeated* in the internal nodes of the B⁺-tree to guide the search


B⁺-tree: Internal node structure

- Each internal node is of the form $\langle P_1, K_1, P_2, K_2, \dots, P_{q-1}, K_{q-1}, P_q \rangle$ where $q \leq p$ and each P_i is a **tree pointer**
- Within each internal node, $K_1 < K_2 < \dots < K_{q-1}$
- For all search field values X in the subtree pointed at by P_i , we have $K_{i-1} < X \leq K_i$ for $1 < i < q$; $X \leq K_i$ for $i = 1$; and $K_{i-1} < X$ for $i = q$



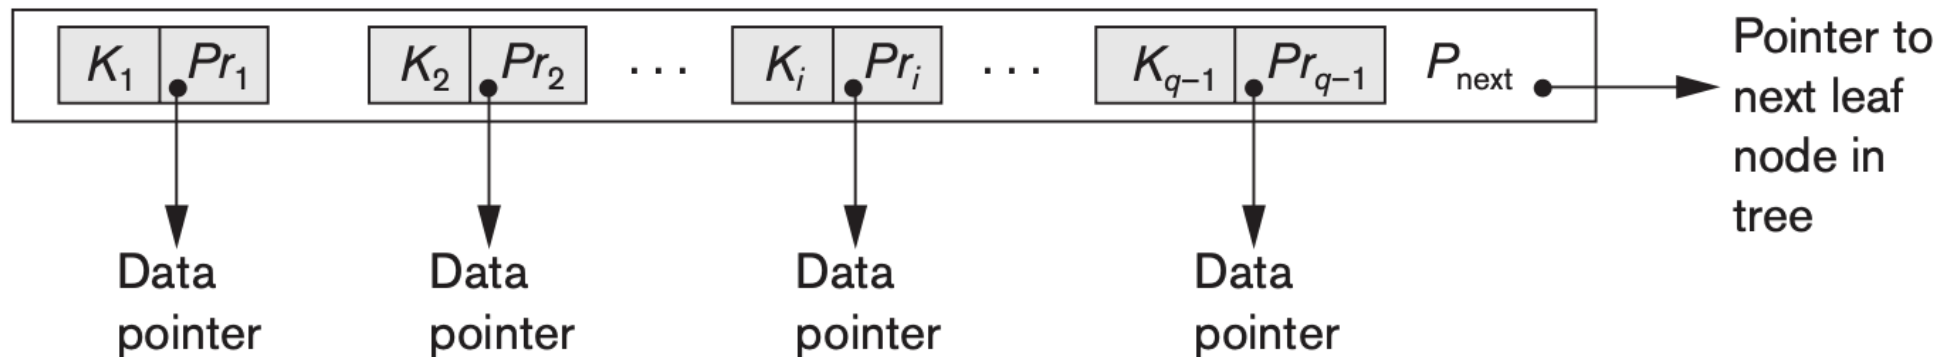
B⁺-tree: Internal node structure (Contd.)

- Each internal node has at most p tree pointers
- Each internal node, except the root, has at least $\lceil (p/2) \rceil$ tree pointers
- The root node has at least two tree pointers if it is an internal node
- An internal node with q pointers, $q \leq p$, has $q - 1$ search field values

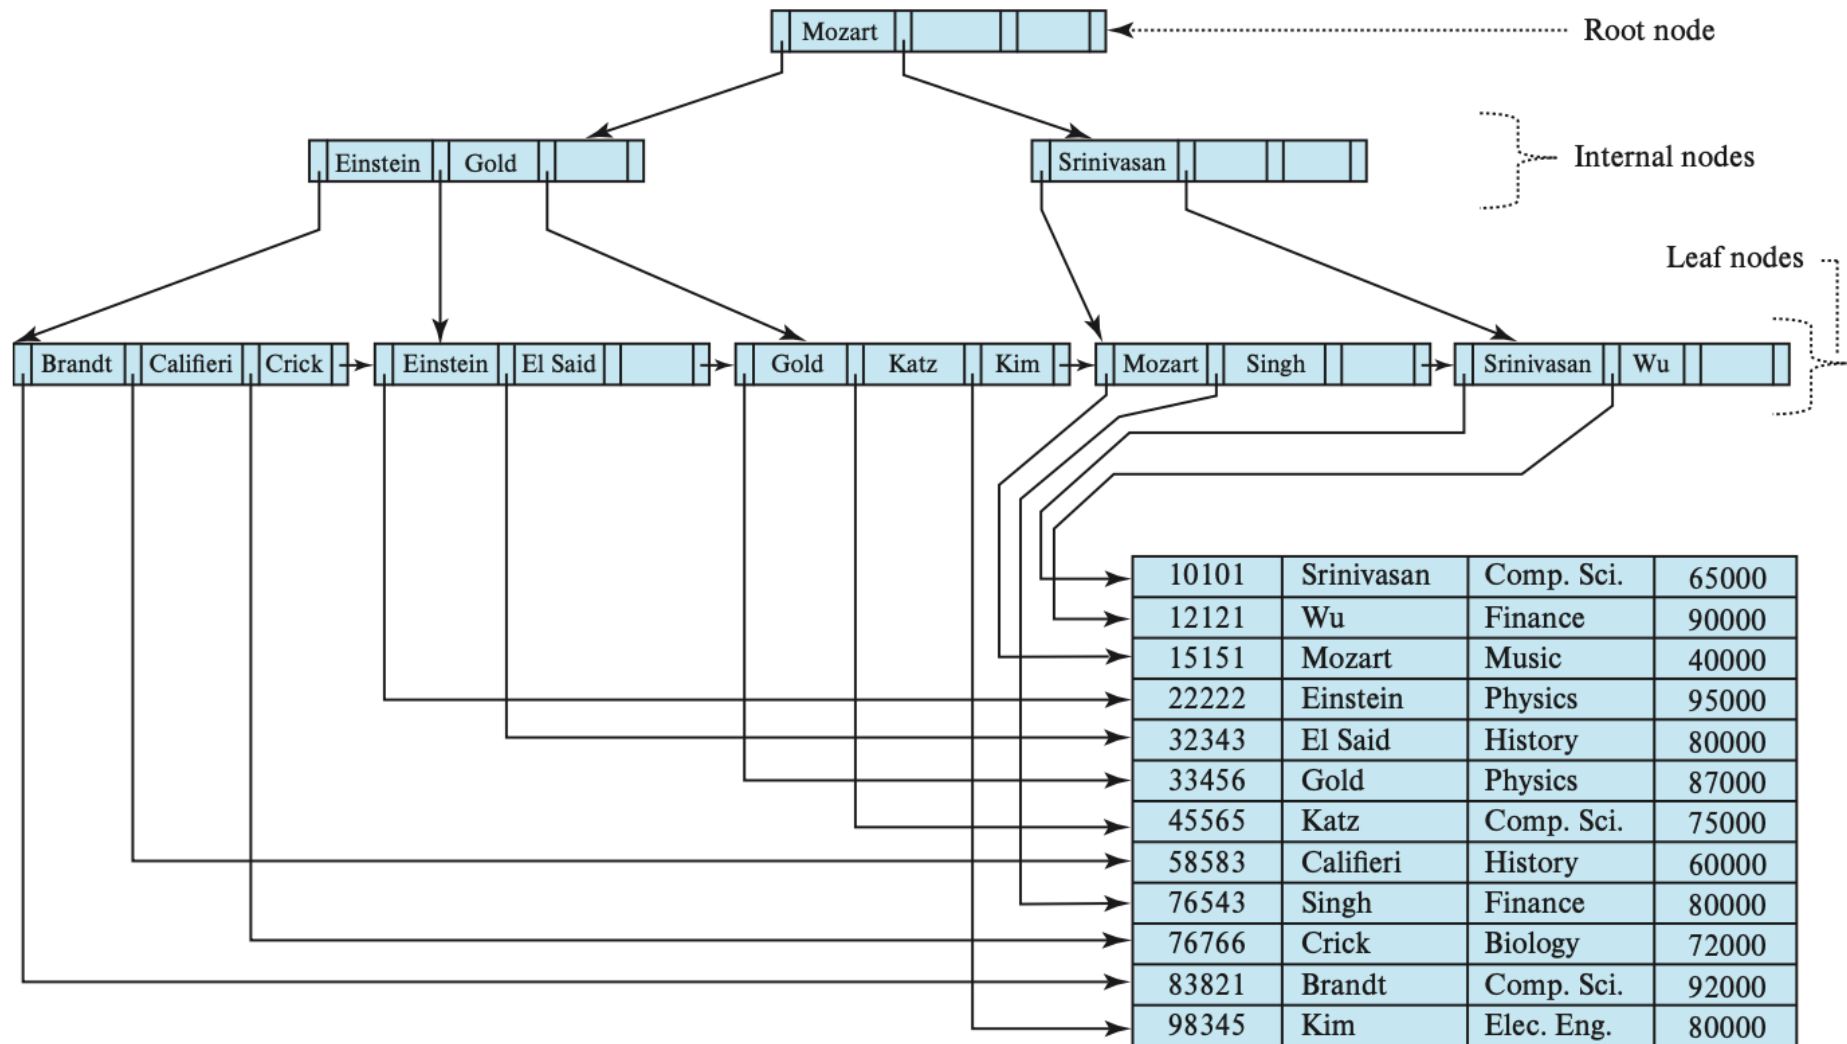


B⁺-tree: Leaf node structure

- Each leaf node is of the form $\langle \langle K_1, Pr_1 \rangle, \langle K_2, Pr_2 \rangle, \dots, \langle K_{q-1}, Pr_{q-1} \rangle, P_{\text{next}} \rangle$ where $q \leq p$, each Pr_i is a data pointer, and P_{next} points to the next *leaf node*
- Within each leaf node, $K_1 \leq K_2 \dots, K_{q-1}$
- Each leaf node has at least $\lceil (p/2) \rceil$ values
- All leaf nodes are at the same level



B⁺-tree: Example



B⁺-tree: Order of nodes

- Because the structures for internal and for leaf nodes of a B+-tree are different, the order p can be different
- Suppose that the search key field is $K = 9$ bytes long, the block size is $B = 512$ bytes, a record pointer is $Pr = 7$ bytes, and a tree pointer is $P = 6$ bytes

$$\begin{aligned}(p * P) + ((p - 1) * K) &\leq B \\(p * 6) + ((p - 1) * 9) &\leq 512 \\(15 * p) &\leq 521; p \leq 34\end{aligned}$$

$$\begin{aligned}(p_{\text{leaf}} * (Pr + K)) + P &\leq B \\(p_{\text{leaf}} * (7 + 9)) + 6 &\leq 512 \\(16 * p_{\text{leaf}}) &\leq 506; p_{\text{leaf}} \leq 31\end{aligned}$$

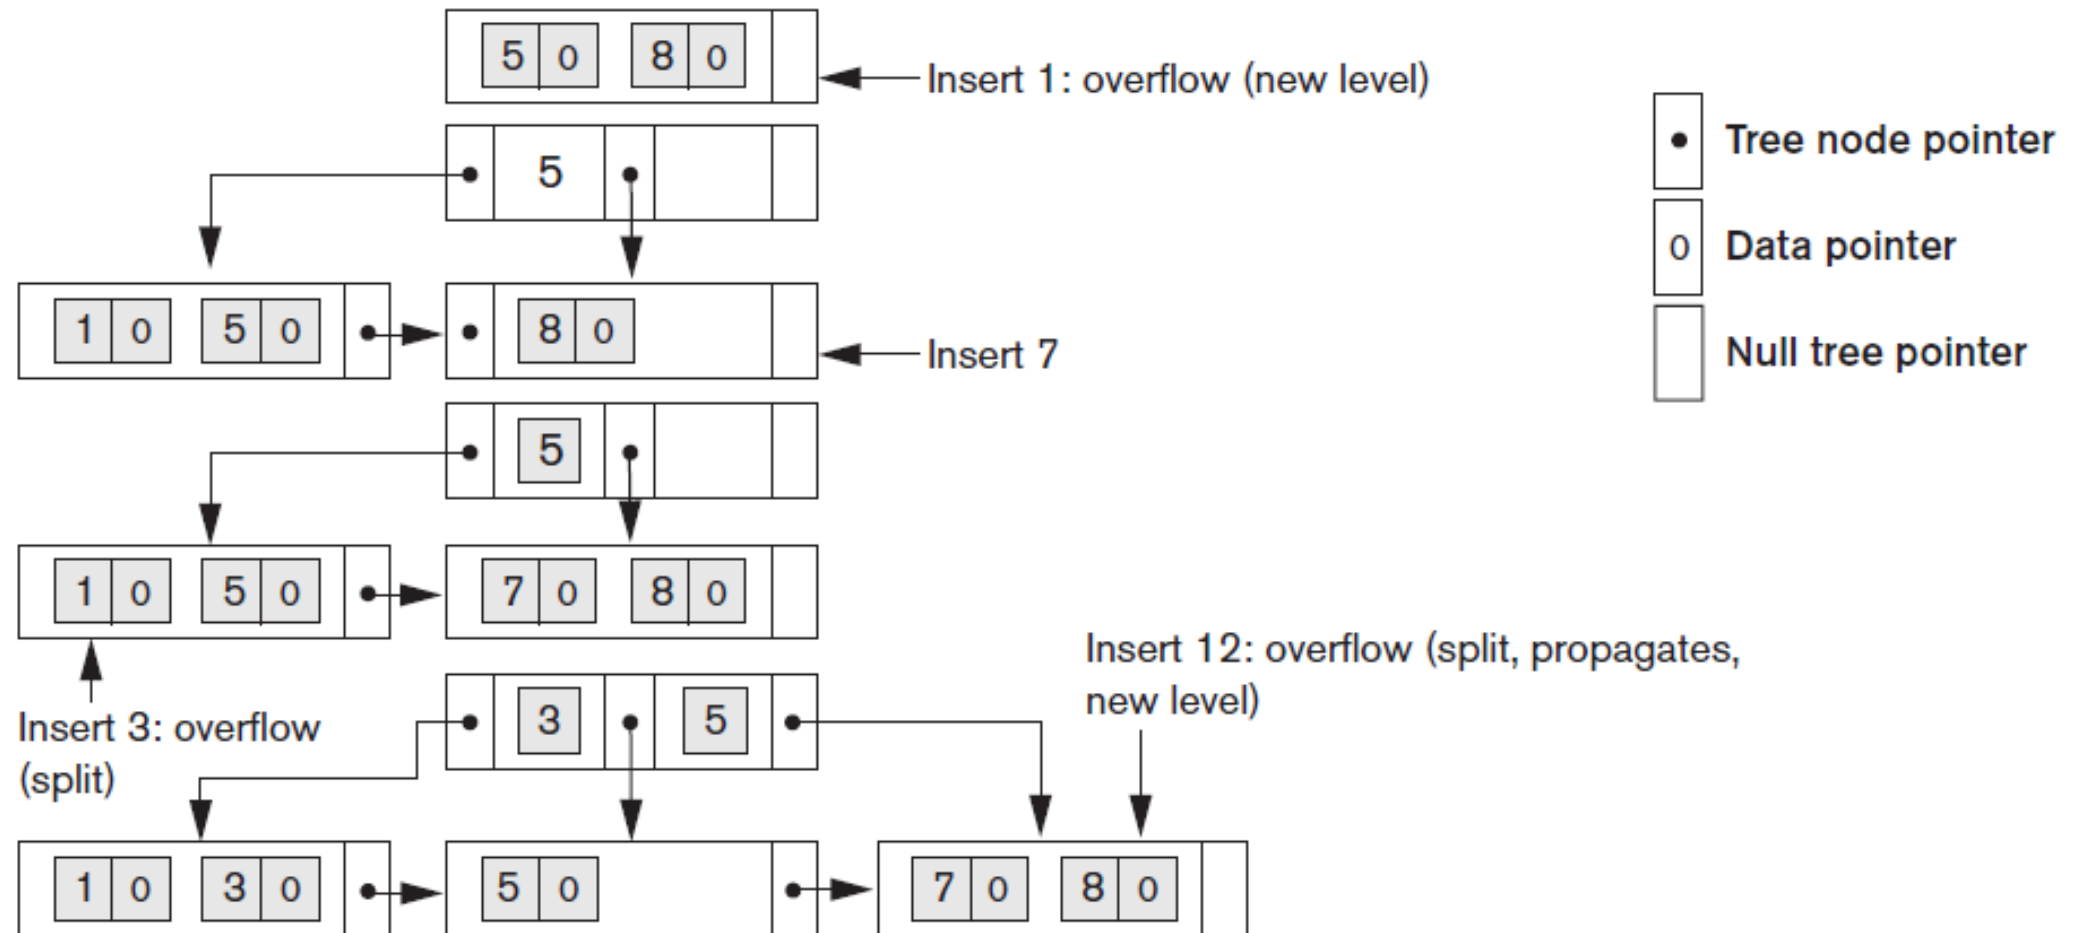
$$\begin{aligned}(p * P) + ((p - 1) * K) + (p - 1) * Pr &\leq B \\(p * 6) + ((p - 1) * 16) &\leq 512 \\(22 * p) &\leq 528; p \leq 24\end{aligned}$$

B⁺-tree: Insertion

- **Splitting a leaf node because of overflow**
 - When a leaf node is full and a new value is inserted there, the node overflows and must be split
 - The first $j = \lceil (p_{leaf} + 1) / 2 \rceil$ entries are kept in the original node, and the remaining entries are moved to a new leaf node
 - The j^{th} value is replicated in the parent internal node
 - If the parent internal is full, the new value (the j^{th} value in above step) will cause it to overflow also, so it must be split

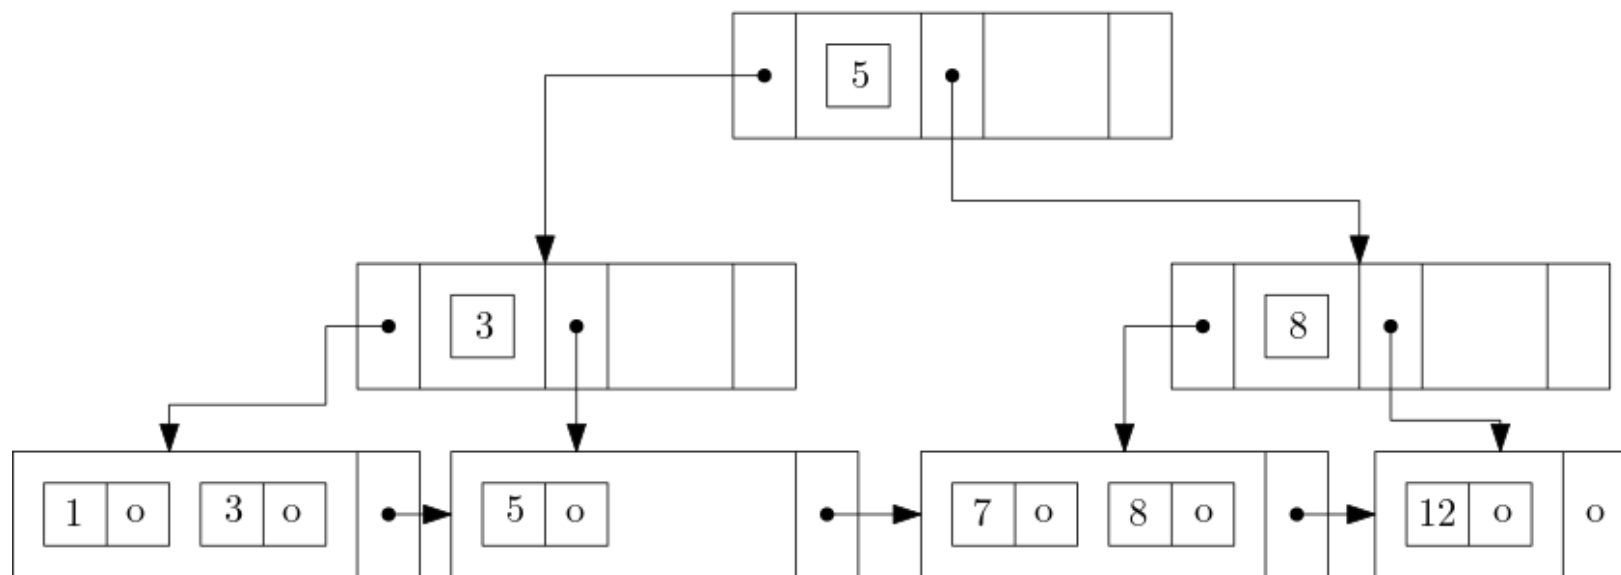


An example of insertion in a B⁺-tree with $p = 3$ and $p_{\text{leaf}} = 2$.

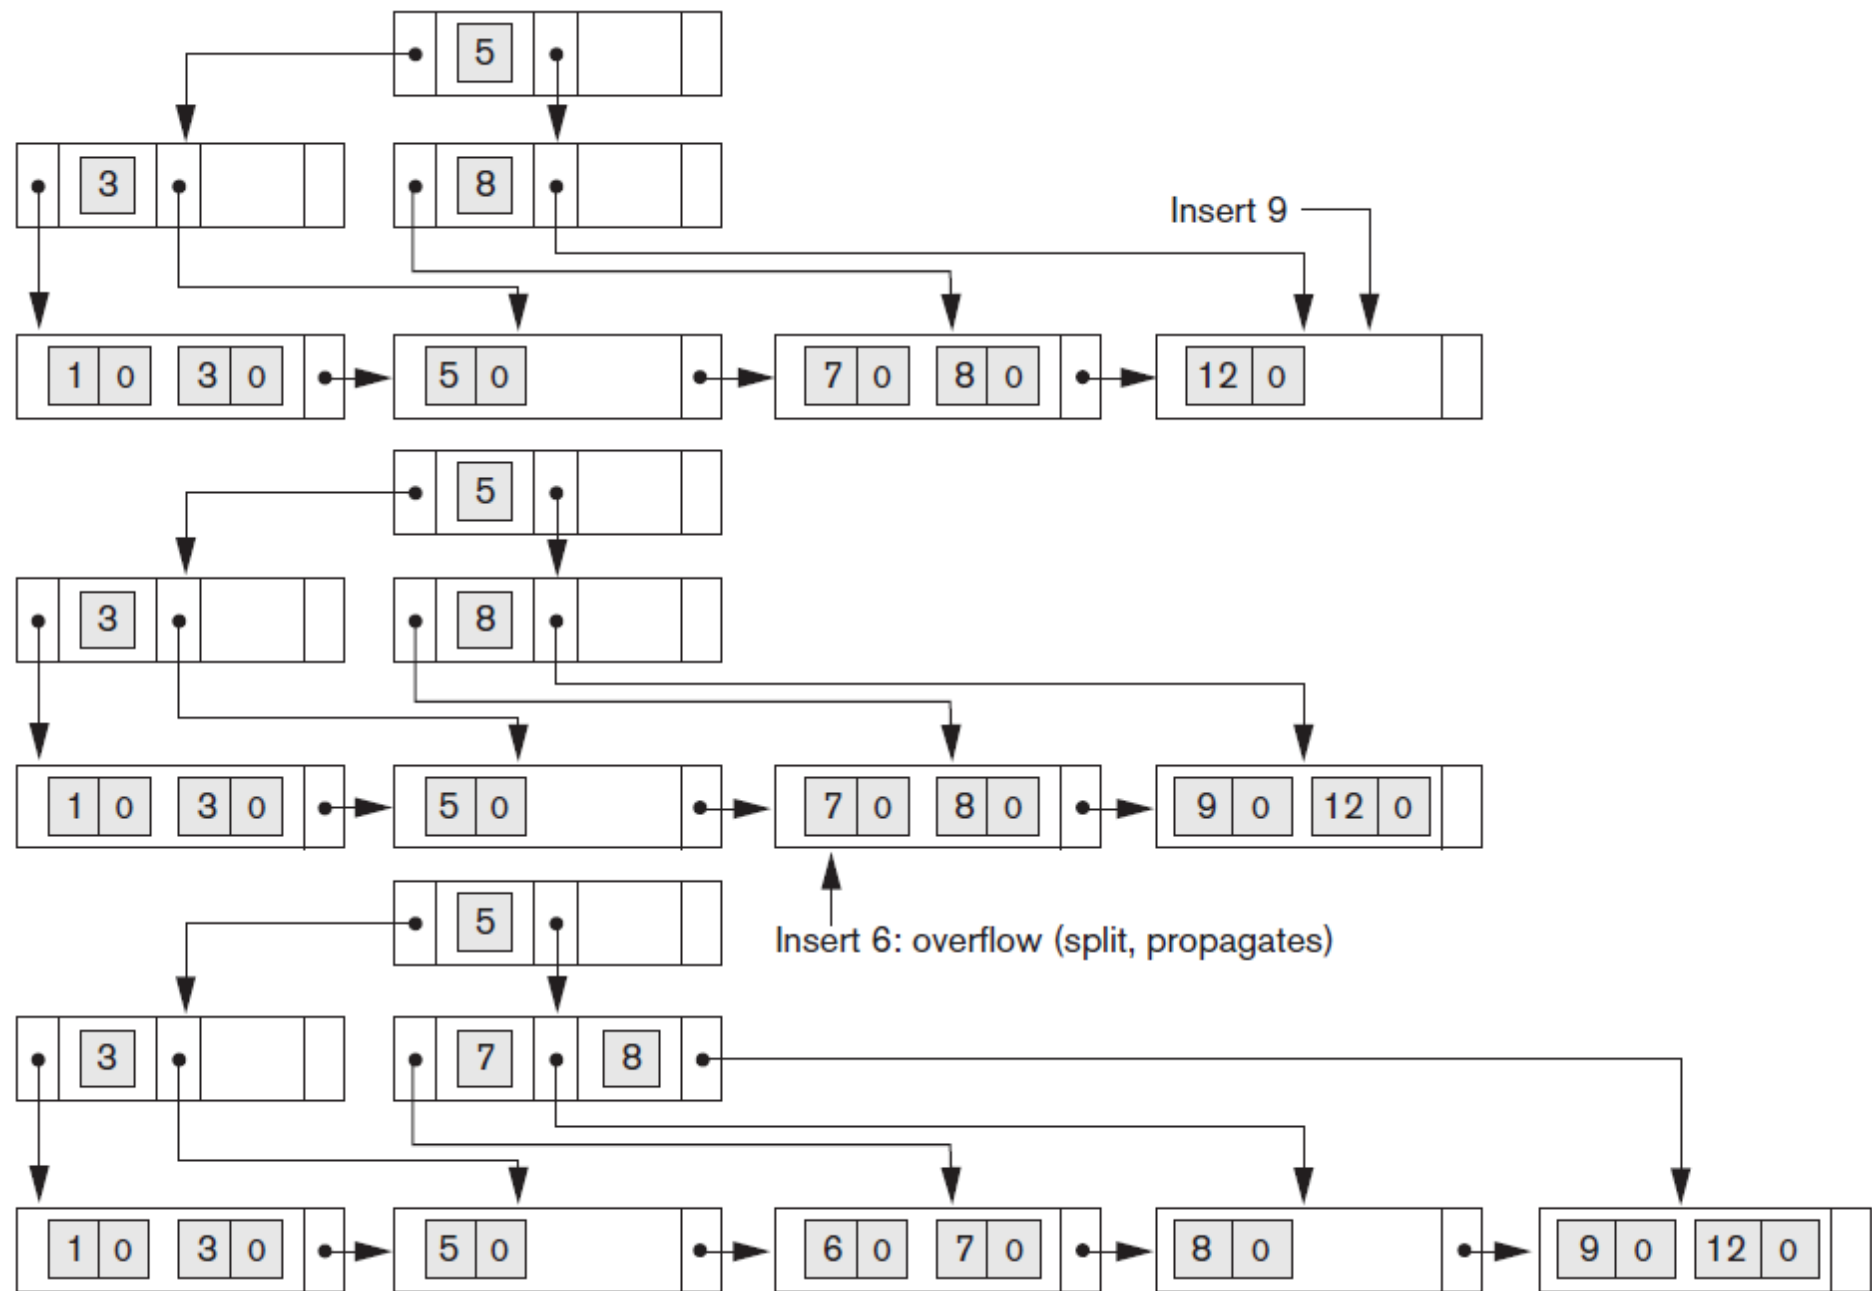
Insertion sequence: 8, 5, 1, 7, 3, 12, 9, 6



B⁺-tree: Insertion (Contd.)

- **Splitting an internal node because of overflow**
 - The entries in the internal node up to P_j , where $j = \lfloor (p + 1)/2 \rfloor$ are kept, while the j^{th} search value is moved to the parent, but not replicated
 - A new internal node will hold the entries from P_{j+1} to the end of the entries in the node
- This splitting can propagate all the way up to create a new root and hence a new level for the B⁺-Tree





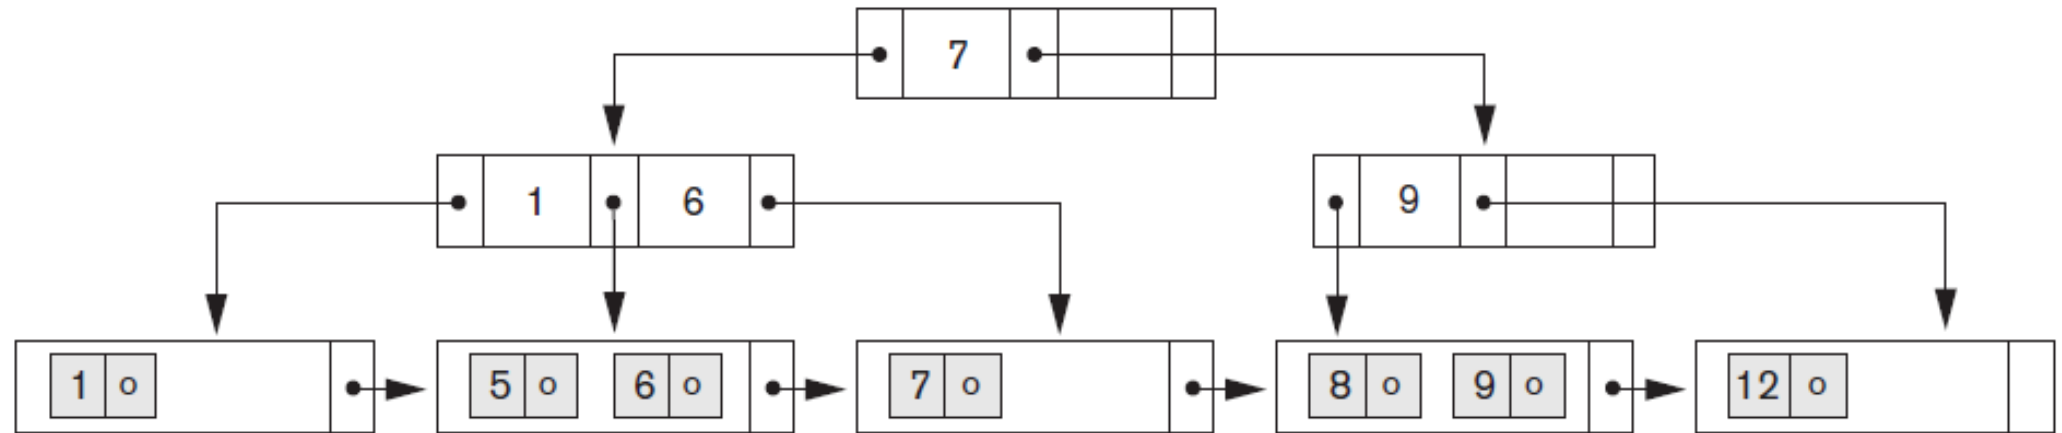
B⁺-tree: Deletion

- When an entry is deleted, it is always removed from the leaf level
- If it happens to occur in an internal node, it must also be removed from there
 - In this case, the value to its left in the leaf node must replace it in the internal node
 - **Reason:** that value is now the rightmost entry in the subtree

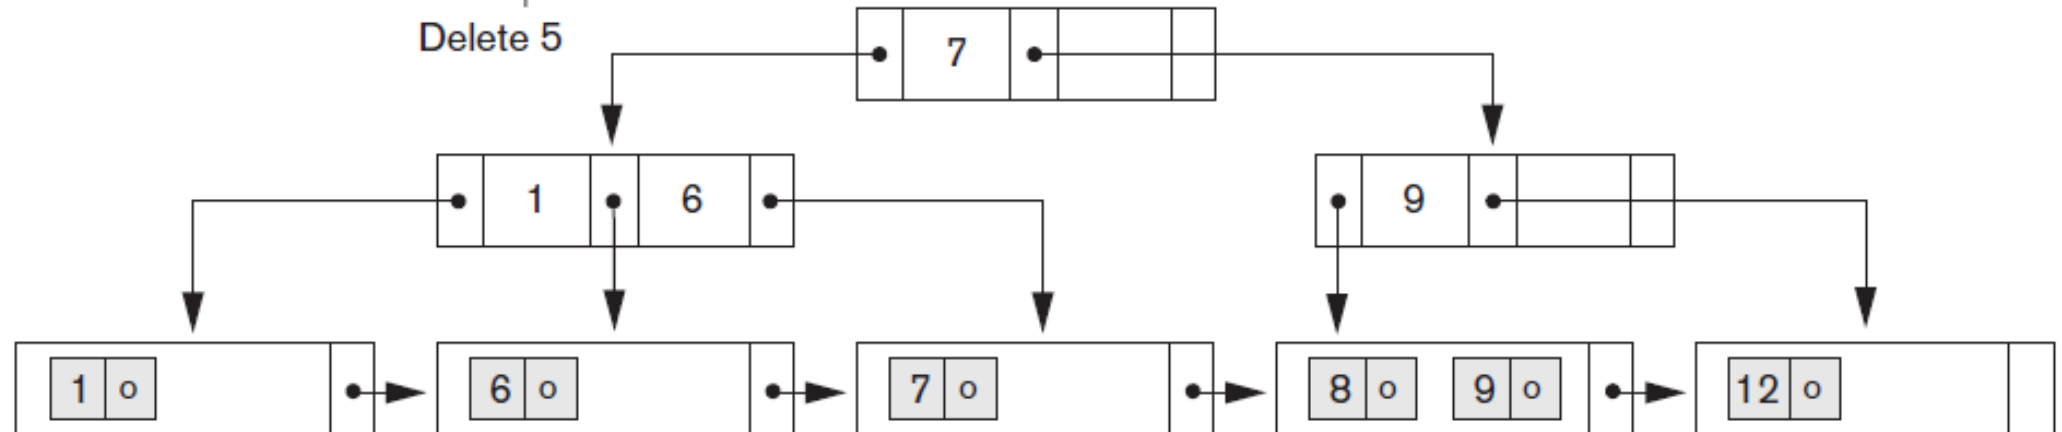
B⁺-tree: Deletion (Contd.)

- **Underflow in the leaf node:**
 - Try to redistribute the entries from the *left or right* sibling node, if possible
 - If the redistribution is not possible, then the three nodes are merged into two leaf nodes
 - In this case, underflow may propagate to internal nodes because one fewer tree pointer and search value are needed
 - This can propagate and reduce the tree levels

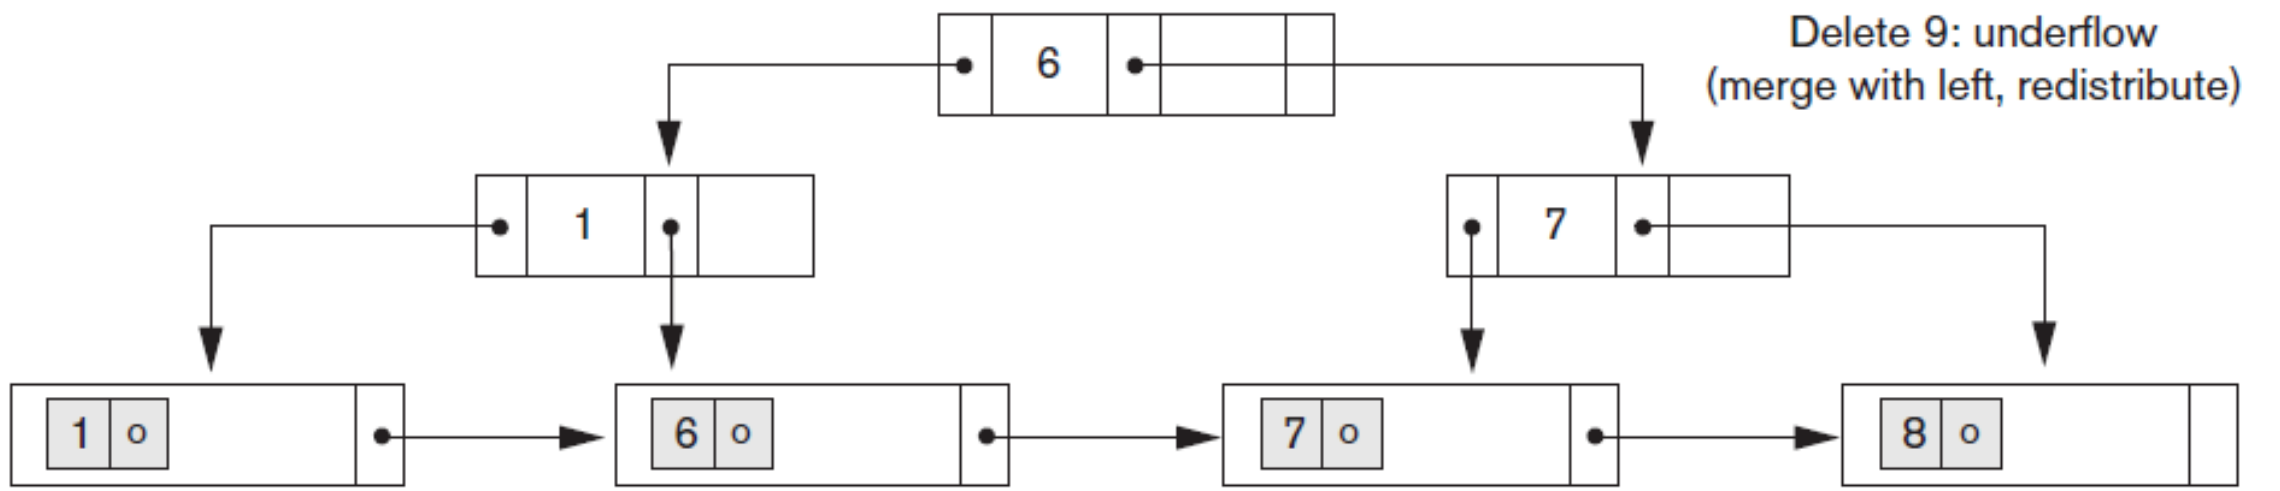
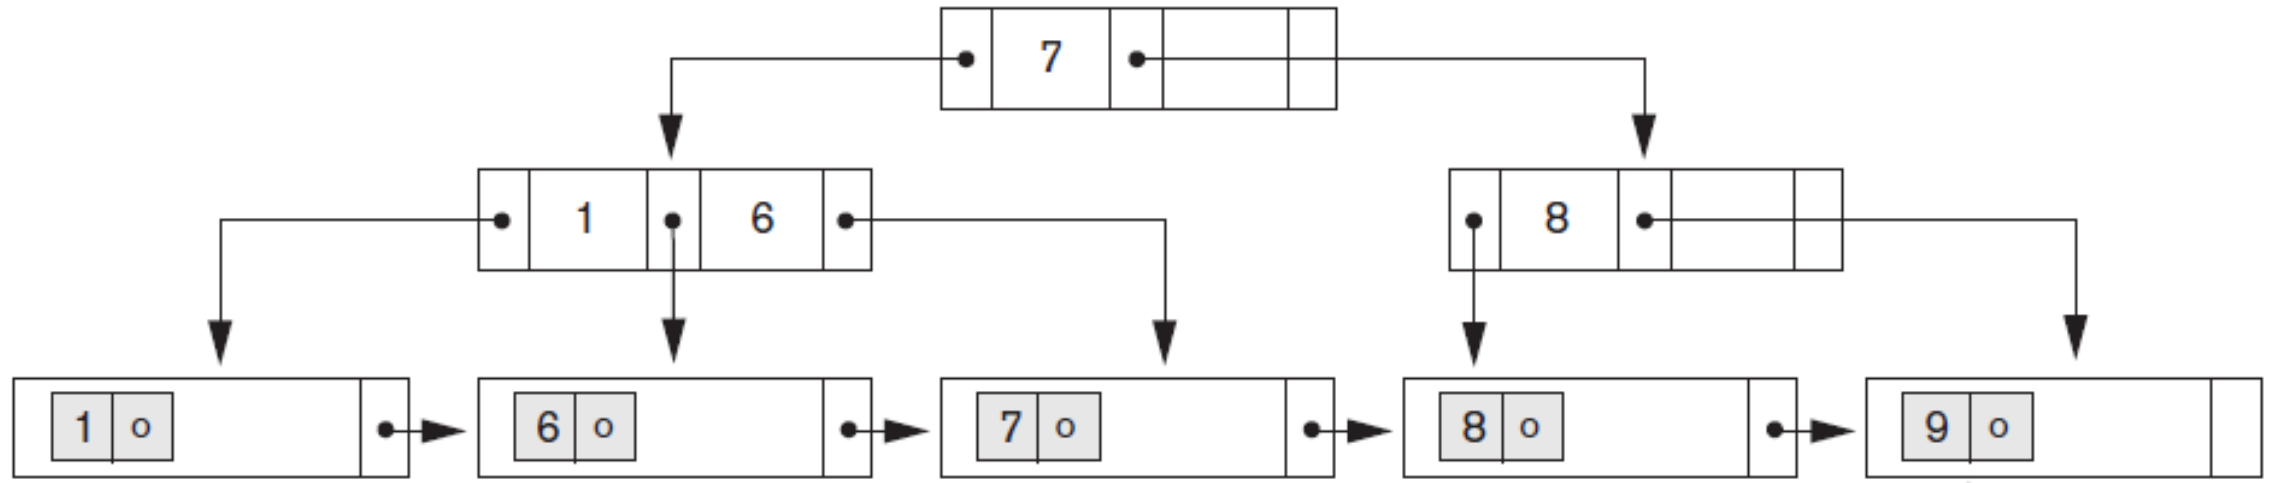
Deletion sequence: 5, 12, 9



Delete 5



Delete 12: underflow
(redistribute)



Thank you!