

# Database Management Systems (DBMS)

Lec 24: Transaction Processing, Concurrency Control, and Recovery

Ramesh K. Jallu  
IIIT Raichur

Date: 08/06/21

# Transactions

- A *transaction* is a collection of operations that form a single logical unit of database processing
  - E.x., transfer of money from one account to another is a transaction, booking a reservation, online purchase, etc.
- A transaction is typically implemented by a computer program that includes database commands such as retrievals, insertions, deletions, and updates
- The operations in a transaction can either be embedded within an application program or they can be specified interactively via SQL

# Transactions (Contd.)

- We focus on the basic concepts and theory that are needed to ensure the correct executions of transactions
- The desirable properties of transactions:
  - *Atomicity*, *Consistency*, *Isolation*, and *Durability*

# ACID properties

- **Atomicity**: Either all operations of the transaction are reflected properly in the database, or none are
- **Consistency**: The integrity constraints must be maintained so that the database is consistent before and after the transaction
  - Ensuring consistency for an individual transaction is the responsibility of the application programmer who codes the transaction
- **Isolation**: Each transaction is unaware of other transactions executing concurrently in the system
- **Durability**: After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures

# Transactions (Contd.)

- A transaction is delimited by statements (or function calls) of the form *begin transaction* and *end transaction*
- The transaction consists of *all operations* executed between the begin transaction and end transaction
- If a transaction *do not update the database but only retrieve data*, the transaction is called a *read-only transaction*; otherwise it is known as a *read-write transaction*

# Transaction model

- We consider a simple database language that focuses on when data are moved from disk to main memory and from main memory to disk
- The only actual operations on the data are restricted in our simple language to arithmetic operations
- A **database** is basically represented as a collection of *named data items* and the size of a data item is called its *granularity*
- The transaction processing concepts we discuss are independent of the data item granularity and apply to data items in general
- The data items in our simplified model have *unique* name and contain a *single* data value

# The two basic operations

- Let us consider a simple bank application consisting of several accounts and a set of transactions that access and update those accounts
- Transactions access data using two operations:
  1. *read\_item(X)*, or simply *read(X)*:
    - Reads a database item named  $X$  into a program variable, also called  $X$ 
      - *I.e.,  $X \leftarrow \text{read}(X)$*
  2. *write\_item(X)*, or simply *write(X)*:
    - Writes the value of program variable  $X$  into the database item named  $X$ 
      - *I.e.,  $X \leftarrow \text{write}(X)$*

# Data transfer from disk to main memory

- The basic unit of data transfer from disk to main memory is one disk page (disk block)
- The DB system maintains a *database cache*, a number of **data buffers** in main memory
- Each buffer typically holds the contents of one database disk block
- When a buffer overflow occurs, buffer replacement polycies are used (such as LRU, FIFO, etc.)
- If the chosen buffer has been modified, it must be written back to disk before it is reused



# The execution of read\_item( $X$ )

1. Find the address of the disk block that contains item  $X$
2. Copy that disk block into a buffer in main memory (if that disk block is not already in some main memory buffer)
  - The size of the buffer is the same as the disk block size
3. Copy item  $X$  from the buffer to the program variable named  $X$

# The execution of write\_item(*X*)

1. Find the address of the disk block that contains item *X*
2. Copy that disk block into a buffer in main memory (if that disk block is not already in some main memory buffer)
3. Copy item *X* from the program variable named *X* into its correct location in the buffer
4. Store the updated disk block from the buffer back to disk (either immediately or at some later point in time)

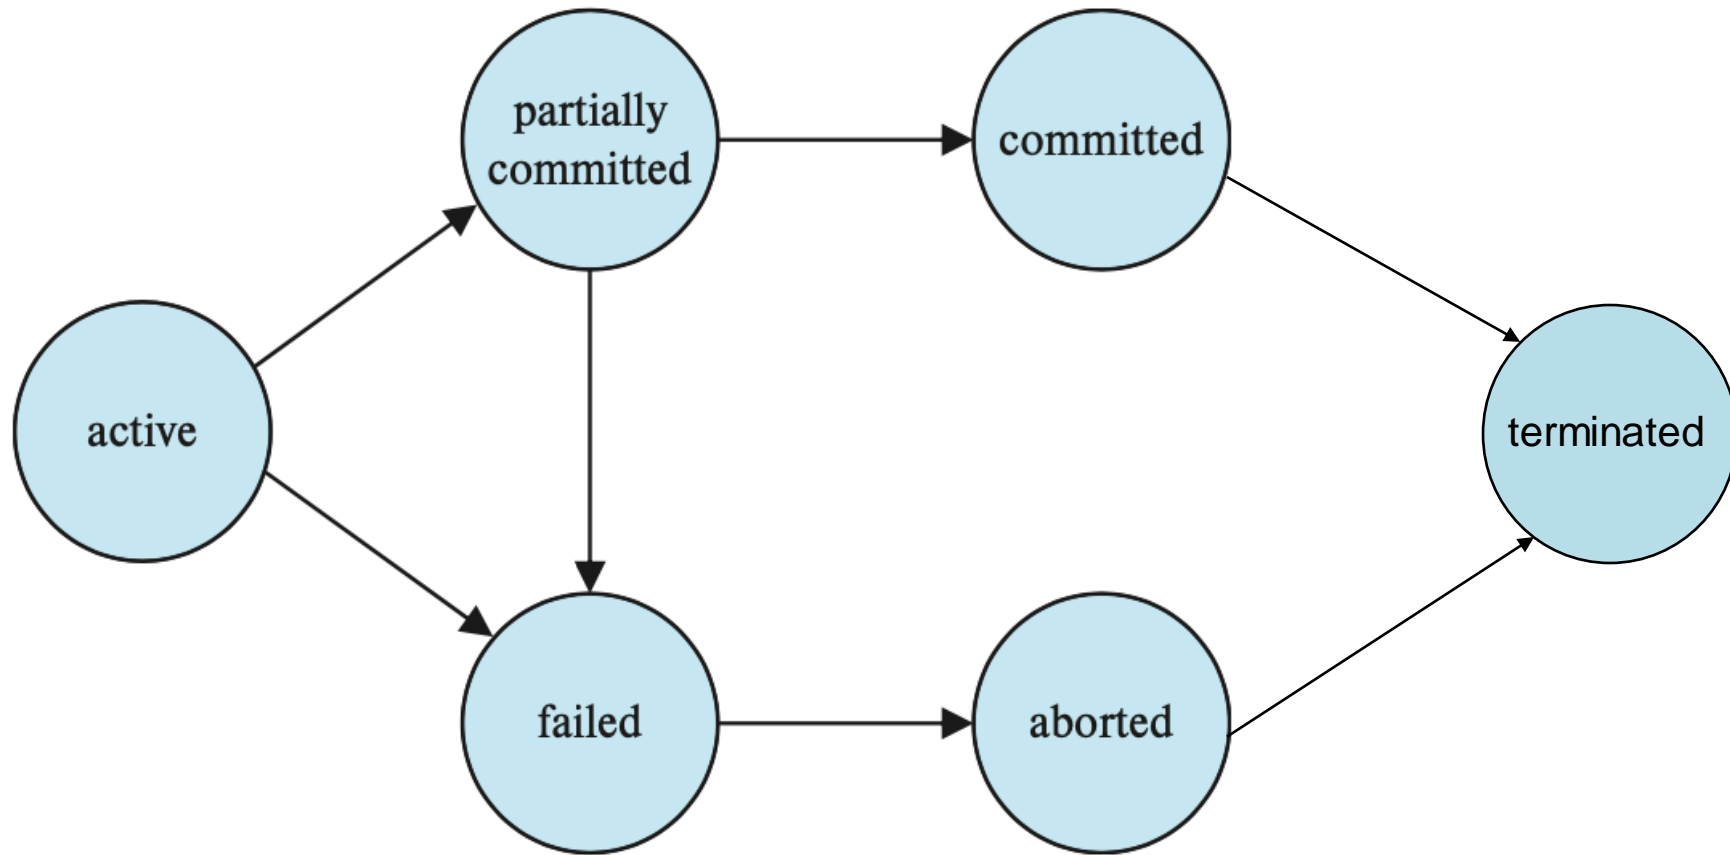
# Example

- Let  $T_i$  be a transaction that transfers Rs. 50 from account  $A$  to account  $B$
- This transaction can be defined as:

$T_i$ : read( $A$ );  
     $A := A - 50$ ;  
    write( $A$ );  
    read( $B$ );  
     $B := B + 50$ ;  
    write( $B$ ).

1. Consistency
2. Atomicity
3. Durable
4. Isolation

# Different states in transactions



# Schedules

- The execution sequences are called *schedules*
- They represent the chronological order in which instructions are executed in the system
- A schedule is a *serial schedule* if each transaction starts only after the previous one has completed
  - **All serial schedules leave the system in consistent state**
- For a set of  $n$  transactions, there exist  $n!$  different valid serial schedules
- Transactions can either be executed serially or concurrently

# Concurrent transactions

- When the database system executes several transactions concurrently, the corresponding schedule no longer needs to be serial
- If two transactions are running concurrently, the operating system may execute one transaction for a little while, then 2nd, then 1st, and so on
- With multiple transactions, the CPU time is shared among all the transactions
- Several execution sequences are possible, since the various instructions from multiple transactions are interleaved

# Concurrent transactions (Contd.)

- Two good reasons for allowing concurrency
  - Improved throughput and resource utilization
  - Reduced waiting time
- Allowing multiple transactions to update data concurrently causes several complications with consistency of the data

**Not all concurrent schedules result in a consistent state**
- When several transactions run concurrently, the isolation property may be violated
- The database system must control the interaction among the concurrent transactions to prevent them from destroying the consistency in the system

# Example: Serial execution

$T_1$ : read( $A$ );  
   $A := A - 50$ ;  
  write( $A$ );  
  read( $B$ );  
   $B := B + 50$ ;  
  write( $B$ ).  
  
 $T_2$ : read( $A$ );  
   $temp := A * 0.1$ ;  
   $A := A - temp$ ;  
  write( $A$ );  
  read( $B$ );  
   $B := B + temp$ ;  
  write( $B$ ).

$T_1$	$T_2$
read( $A$ ) $A := A - 50$ write( $A$ ) read( $B$ ) $B := B + 50$ write( $B$ ) commit	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ ) read( $B$ ) $B := B + temp$ write( $B$ ) commit

$T_1$	$T_2$
read( $A$ ) $A := A - 50$ write( $A$ ) read( $B$ ) $B := B + 50$ write( $B$ ) commit	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ ) read( $B$ ) $B := B + temp$ write( $B$ ) commit



# Example: Concurrent execution

$T_1$ : read( $A$ );  
 $A := A - 50$ ;  
 write( $A$ );  
 read( $B$ );  
 $B := B + 50$ ;  
 write( $B$ ).  
  
 $T_2$ : read( $A$ );  
 $temp := A * 0.1$ ;  
 $A := A - temp$ ;  
 write( $A$ );  
 read( $B$ );  
 $B := B + temp$ ;  
 write( $B$ ).

$T_1$	$T_2$
read( $A$ ) $A := A - 50$ write( $A$ )	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ )
read( $B$ ) $B := B + 50$ write( $B$ ) commit	read( $B$ ) $B := B + temp$ write( $B$ ) commit


$T_1$	$T_2$
read( $A$ ) $A := A - 50$	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ ) read( $B$ )
write( $A$ ) read( $B$ ) $B := B + 50$ write( $B$ ) commit	$B := B + temp$ write( $B$ ) commit

# Serializability

- We can ensure the consistency under concurrent execution by making sure that the concurrent execution has the same effect as a serial schedule
  - That is, the schedule should be equivalent to a serial schedule
- Such schedules are called *serializable* schedules
- How to determine a schedule is serializable?

# Conflict serializability

- Let us consider a schedule  $S$  in which there are two consecutive instructions,  $I$  and  $J$ , of transactions  $T_i$  and  $T_j$ , respectively ( $i \neq j$ )
- If  $I$  and  $J$  refer to different data items, then we can swap  $I$  and  $J$  without affecting the results of any instruction in the schedule
- If  $I$  and  $J$  refer to the same data item  $X$ , then the order of the two steps may matter: there are four cases that we need to consider
  - $I = \text{read}(X), J = \text{read}(X)$ ;
  - $I = \text{read}(X), J = \text{write}(X)$ ;
  - $I = \text{write}(X), J = \text{read}(X)$ ;
  - $I = \text{write}(X), J = \text{write}(X)$ ;



# Conflict serializability (Contd.)

- We say that  $I$  and  $J$  *conflict* if they are operations by different transactions on the same data item, and at least one of these instructions is a *write* operation
- If  $I$  and  $J$  do not conflict, then we can swap the order of  $I$  and  $J$  to produce a new schedule  $S'$
- $S$  is equivalent to  $S'$ , since all instructions appear in the same order in both schedules except for  $I$  and  $J$ , whose order does not matter
- If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of nonconflicting instructions, we say that  $S$  and  $S'$  are *conflict equivalent*

# Conflict serializability (Contd.)

$T_1$	$T_2$
read( $A$ ) write( $A$ )	read( $A$ ) write( $A$ )
read( $B$ ) write( $B$ )	
	read( $B$ ) write( $B$ )

$T_1$	$T_2$
read( $A$ ) write( $A$ )	read( $A$ )
read( $B$ )	
write( $B$ )	write( $A$ )
	read( $B$ ) write( $B$ )

$T_1$	$T_2$
read( $A$ ) write( $A$ ) read( $B$ ) write( $B$ )	read( $A$ ) write( $A$ ) read( $B$ ) write( $B$ )

# Conflict serializability (Contd.)

- Are serial schedules conflict equivalent to each other?
  - I.e., if  $S$  and  $T$  are two serial schedules, can we obtain  $S$  from  $T$  (or vice versa) by a series of swapings?
- We say that a schedule  $S$  is *conflict serializable* if it is conflict equivalent to a serial schedule

$T_3$	$T_4$
read( $Q$ )	write( $Q$ )
write( $Q$ )	

Thank you!