

# Database Management Systems (DBMS)

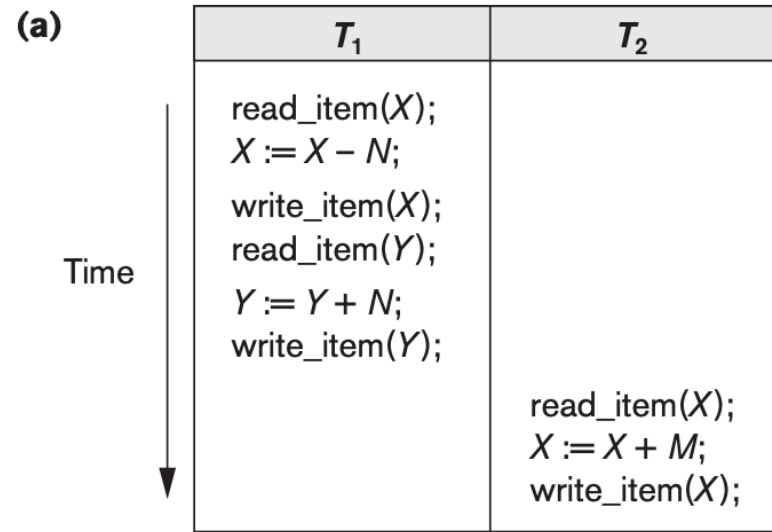
Lec 25: Transaction Processing, Concurrency Control,  
and Recovery

Ramesh K. Jallu  
IIIT Raichur

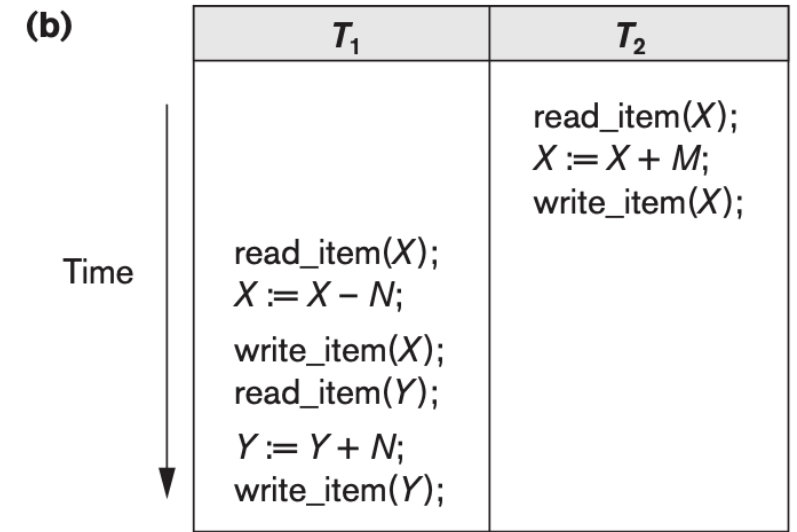
Date: 11/06/21

# Recap

- Desirable properties
  - **A**tomicity
  - **C**onsistency
  - **I**solation
  - **D**urability
- States in transaction
- **Schedules**: the order of execution of operations from all the various transactions
  - **Serial schedule**: All serial schedules leave the system in consistence state
  - **Concurrent execution**: Not all leave the system in consistence state
- **Conflict serializability**: Conflict equivalent to a serial schedule

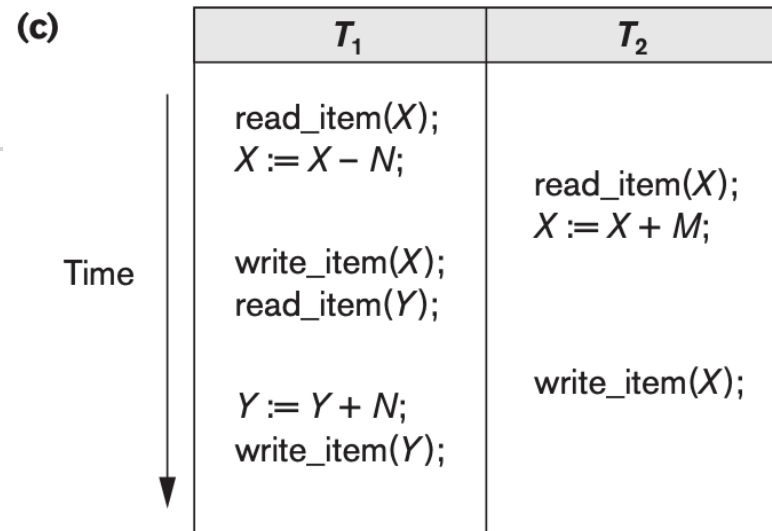


Schedule A

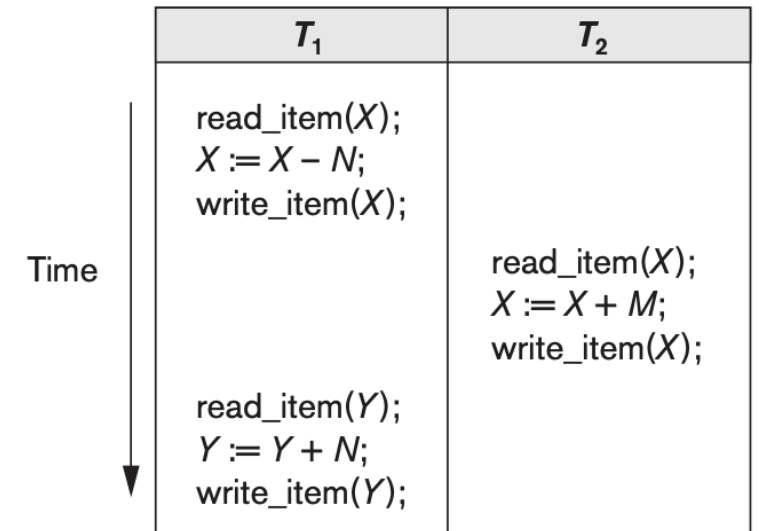


Schedule B

# Example



Schedule C




Schedule D

# Determining a conflict serializability

- Let  $S$  be a schedule for transactions  $T_1, T_2, \dots, T_n$
- A *precedence graph* (or *serialization graph*) for  $S$  is a directed graph  $G = (V, E)$ , where  $V = \{T_1, T_2, \dots, T_n\}$ , and there is an edge between two vertices  $T_i$  and  $T_j$  if one of the following conditions holds
  1.  $T_i$  executes *write*( $X$ ) before  $T_j$  executes *read*( $X$ )
  2.  $T_i$  executes *read*( $X$ ) before  $T_j$  executes *write*( $X$ )
  3.  $T_i$  executes *write*( $X$ ) before  $T_j$  executes *write*( $X$ )


# Example-1

	Transaction $T_1$	Transaction $T_2$	Transaction $T_3$
Time ↓	read_item(X); write_item(X);	read_item(Z); read_item(Y); write_item(Y);	read_item(Y); read_item(Z);  write_item(Y); write_item(Z);
	read_item(Y); write_item(Y);	read_item(X);  write_item(X);	




# Example-2

Transaction $T_1$	Transaction $T_2$	Transaction $T_3$
<div data-bbox="50 678 191 1192" data-label="Text">Time ↓</div> <div data-bbox="267 721 535 821">read_item(<math>X</math>); write_item(<math>X</math>);</div> <div data-bbox="267 985 535 1085">read_item(<math>Y</math>); write_item(<math>Y</math>);</div>	<div data-bbox="751 928 1006 978">read_item(<math>Z</math>);</div> <div data-bbox="751 1056 1019 1249">read_item(<math>Y</math>); write_item(<math>Y</math>); read_item(<math>X</math>); write_item(<math>X</math>);</div>	<div data-bbox="1197 642 1465 742">read_item(<math>Y</math>); read_item(<math>Z</math>);</div> <div data-bbox="1197 821 1465 921">write_item(<math>Y</math>); write_item(<math>Z</math>);</div>



# Observations

- If the precedence graph for  $S$  has a cycle, then schedule  $S$  is not conflict serializable; otherwise,  $S$  is conflict serializable
- If the graph is a DAG, a *serializable order* can be found by topological sorting
- In general, several possible linear orders that can be obtained through a topological sort



$\langle T_i, T_j, T_k, T_m \rangle$

$\langle T_i, T_k, T_j, T_m \rangle$

# Why is concurrency control is needed?

- When several transactions execute concurrently in the database, the consistency of data may no longer be preserved
- Several problems can occur when concurrent transactions execute in an uncontrolled manner
- It is necessary for the system to control the interaction among the concurrent transactions, and this control is achieved through one of a variety of mechanisms called *concurrency control* schemes.



# Example

- Consider an airline reservations database in which a record is stored for each airline flight
- Let  $T_1$  transfers  $N$  reservations from one flight whose number of reserved seats is stored in the database item named  $X$  to another flight whose number of reserved seats is stored in the database item named  $Y$
- Let  $T_2$  that just reserves  $M$  seats on the first flight ( $X$ ) referenced in transaction  $T_1$

$T_1$
read_item( $X$ ); $X := X - N$ ; write_item( $X$ ); read_item( $Y$ ); $Y := Y + N$ ; write_item( $Y$ );

$T_2$
read_item( $X$ ); $X := X + M$ ; write_item( $X$ );

# 1. The lost update problem

$T_1$
read_item(X); $X := X - N$ ; write_item(X); read_item(Y); $Y := Y + N$ ; write_item(Y);

$T_2$
read_item(X); $X := X + M$ ; write_item(X);

$T_1$	$T_2$
read_item(X); $X := X - N$ ;  write_item(X); read_item(Y);  $Y := Y + N$ ; write_item(Y);	  read_item(X); $X := X + M$ ;   write_item(X);

## 2. The temporary update (or dirty read) problem

$T_1$
read_item(X); $X := X - N$ ; write_item(X); read_item(Y); $Y := Y + N$ ; write_item(Y);

$T_2$
read_item(X); $X := X + M$ ; write_item(X);

$T_1$	$T_2$
read_item(X); $X := X - N$ ; write_item(X);       read_item(Y);	       read_item(X); $X := X + M$ ; write_item(X);

### 3. The incorrect summary problem

$T_1$
<pre>read_item(X); X := X - N; write_item(X); read_item(Y); Y := Y + N; write_item(Y);</pre>

$T_2$
<pre>read_item(X); X := X + M; write_item(X);</pre>

$T_1$	$T_3$
<pre>read_item(X); X := X - N; write_item(X);</pre> <pre>read_item(Y); Y := Y + N; write_item(Y);</pre>	<pre>sum := 0; read_item(A); sum := sum + A; ⋮</pre> <pre>read_item(X); sum := sum + X; read_item(Y); sum := sum + Y;</pre>

## 4. Unpredictable read problem

- If transaction  $T$  reads the same data item twice and the data item is changed by another transaction  $T'$  between the two reads
- Hence,  $T$  receives *different values* for its two reads of the same item
- For example, if during an airline reservation transaction, a customer inquires about seat availability on several flights, and upon deciding a particular flight it may end up reading a different value for the item
- The ***phantom read problem*** occurs when a transaction reads a variable once but when it tries to read that same variable again, an error occurs saying that the variable does not exist

# Types of failures

1. **A computer failure:** System crash due to h/w or s/w failure
2. **A transaction or system error:** Logical error
3. **Local errors:** Data not found or insufficient account balance
4. **Concurrency control enforcement:** Deadlock or violates serializability
5. **Disk failure:** Disk read/write head crash
6. **Physical problems and catastrophes:** Power or air-conditioning failure, fire, theft, sabotage, etc.

# Characterizing Schedules Based on Recoverability

- For some schedules it is easy to recover from transaction and system failures, whereas for other schedules the recovery process can be quite involved
- In some cases, it is even not possible to recover correctly after a failure
- We theoretically characterize the different types of schedules for which *recovery is possible*, as well as those for which *recovery is relatively simple*

# Characterizing Schedules Based on Recoverability (Contd.)

- Once a transaction is committed, it should *never* be rolled back T; Otherwise, it violates durability property
- The schedules that theoretically meet this criterion are called *recoverable schedules*
- A schedule where a committed transaction may have to be rolled back during recovery is called *nonrecoverable*



# Recoverable schedules

- A schedule  $S$  is recoverable if no transaction  $T$  in  $S$  commits until all transactions  $T'$  that have written some item  $X$  that  $T$  reads have committed
- A transaction  $T$  **reads** from transaction  $T'$  in a schedule  $S$  if
  1. some item  $X$  is first written by  $T'$  and later read by  $T$ ,
  2.  $T'$  should not have been aborted before  $T$  reads item  $X$ , and
  3. there should be no transactions that write  $X$  after  $T'$  writes it and before  $T$  reads it (unless those transactions, if any, have aborted before  $T$  reads  $X$ )

Thank you!