

Database Management Systems (DBMS)

Lec 26: Transaction Processing, Concurrency Control, and Recovery (Contd.)

Ramesh K. Jallu
IIIT Raichur

Date: 15/06/21

Recap

- Determining a conflict serializability through precedence graph
- Need for concurrency control
 - The lost update problem
 - The temporary or dirty read problem
 - The incorrect summary problem
 - The unpredictable read problem
- Characterizing schedules based on recoverability

Recoverable schedules (Recap)

- A transaction T_j that is dependent on T_i (i.e., T_j has read data written by T_i) commits only after T_i is committed
- The atomicity property requires that any transaction T_j that is dependent on T_i is also must be aborted if T_i fails for whatever reason

T_6	T_7
read(A) write(A)	read(A) commit
read(B)	

T_8	T_9	T_{10}
read(A) read(B) write(A)	read(A) write(A)	read(A)
abort		

Cascadeless schedules

- A phenomenon, in which a single transaction failure leads to a series of transaction rollbacks, is called **cascading rollback**
- It is desirable to restrict the schedules to those where cascading rollbacks cannot occur. Such schedules are called *cascadeless* schedules
- Every cascadeless schedule is a recoverable schedule as every transaction reads items that were written by committed transactions

T_8	T_9	T_{10}
read(A) read(B) write(A)	read(A) write(A)	read(A)
abort		

Concurrency control

- When several transactions execute concurrently in the database the isolation property may no longer be preserved
- The system must control the interaction among the concurrent transactions
- The mechanisms that controls the interaction among transactions is called the *concurrency control schemes* (ptotocols)
- The most commonly used protocols in practise
 1. *Lock based protocols*
 2. *Timestamp based protocols*

Lock based protocols

- To ensure isolation is preserved, the data items be accessed in a *mutually exclusive* manner
 - That is, while one transaction is accessing a data item, no other transaction can modify that data item
- A *lock* is a *variable* associated with a data item that describes the status of the item
- The DBMS has a *lock manager subsystem* to keep track of and control access to locks. The system grants a transaction to access a data item only if it is currently holding a *lock* on that item
- Generally, there is one lock for each data item in the database

Types of locks

- Binary locks
- Shared/Exclusive (or Read/Write) Locks

Binary locks

- A **binary lock** can have two *states* or *values*: locked (1) and unlocked (0)
- A distinct lock is associated with each database item X
- If the value of the lock on X is 1, item X "*cannot be accessed*" by a database operation that requests the item
- If the value of the lock on X is 0, the item X "*can be accessed*" when requested, and the lock value is changed to 1
- We refer to the current value (or state) of the lock associated with item X as *lock*(X)

Operations on binary locks

- Two operations are used in binary locks : *lock_item* and *unlock_item*
- A transaction requests access to an item X by first issuing a *lock_item(X)* operation
 - If $lock(X) = 1$, the transaction is forced to wait
 - If $lock(X) = 0$, it is set to 1 (the transaction **locks** the item) and the transaction is allowed to access item X
- After execution/abort, the transaction issues an *unlock_item(X)* operation, which sets $lock(X)$ back to 0 so that X may be accessed by other transactions
- Hence, a binary lock enforces **mutual exclusion** on the data item

Lock and unlock operations

lock_item(X):

```
B:  if LOCK(X) = 0                (*item is unlocked*)
      then LOCK(X) ← 1            (*lock the item*)
    else
      begin
        wait (until LOCK(X) = 0
              and the lock manager wakes up the transaction);
        go to B
      end;
```

unlock_item(X):

```
LOCK(X) ← 0;                (* unlock the item *)
if any transactions are waiting
  then wakeup one of the waiting transactions;
```

Rules associated with binary locks

- Every transaction must obey the following four rules:
 1. A transaction T must issue the operation $lock_item(X)$ before any $read_item(X)$ or $write_item(X)$ operations are performed in T
 2. A transaction T must issue the operation $unlock_item(X)$ after all $read_item(X)$ and $write_item(X)$ operations are completed in T
 3. A transaction T should not issue a $lock_item(X)$ operation if it already holds the lock on item X
 4. A transaction T should not issue an $unlock_item(X)$ operation unless it already holds the lock on item X

Advantages and disadvantages with binary locks

- It is simple to implement a binary lock; all that is needed is a binary-valued variable, *lock* , associated with each data item *X* in the database
- The system needs maintain a lock table such as a hash table to keep track the *only the items that are currently locked* . Items not in the lock table are considered to be unlocked
- The transactions that are waiting to access a locked item can be put in a queue
- **At most one transaction** can hold the lock on a particular item. Thus no two transactions can access the same item concurrently

Shared/Exclusive (or Read/Write) Locks

- If a transaction access X just for *reading purposes only*, then it doesn't need any lock. However, if the transaction *is to write* an item X , it must have exclusive access to X
- For this purpose a *multiple-mode locks* are used and are called *shared/exclusive* or *read/write*
- **Shared:** If a transaction T_i has obtained a **shared-mode lock** on item X , then T_i can read, but cannot write, X
- **Exclusive:** If a transaction T_i has obtained an **exclusive-mode lock** on item X , then T_i can both read and write X

Operations

- Shared/exclusive mode locking scheme has three locking operations on an item X
 - $read_lock(X)$ (or) $shared_lock(X)$
 - $write_lock(X)$ (or) $exclusive_lock(X)$
 - $unlock(X)$
- Every transaction must **request** a lock in an appropriate mode on data item X , depending on the types of operations that it will perform on

Locking and unlocking operations

read_lock(X):

```
B:  if LOCK(X) = "unlocked"
      then begin LOCK(X) ← "read-locked";
           no_of_reads(X) ← 1
      end
    else if LOCK(X) = "read-locked"
      then no_of_reads(X) ← no_of_reads(X) + 1
    else begin
          wait (until LOCK(X) = "unlocked"
                and the lock manager wakes up the transaction);
          go to B
        end;
```

write_lock(X):

```
B:  if LOCK(X) = "unlocked"
      then LOCK(X) ← "write-locked"
    else begin
          wait (until LOCK(X) = "unlocked"
                and the lock manager wakes up the transaction);
          go to B
        end;
```

unlock (X):

```
if LOCK(X) = "write-locked"
  then begin LOCK(X) ← "unlocked";
        wakeup one of the waiting transactions, if any
      end
else if LOCK(X) = "read-locked"
  then begin
        no_of_reads(X) ← no_of_reads(X) - 1;
        if no_of_reads(X) = 0
          then begin LOCK(X) = "unlocked";
                wakeup one of the waiting transactions, if any
            end
      end;
```

Rules enforced in shared/exclusive locking scheme

1. A transaction T must issue the operation $read_lock(X)$ or $write_lock(X)$ before any $read$ operation on X is performed in T
2. A transaction T must issue the operation $write_lock(X)$ before any $write$ operation is performed on X in T
3. A transaction T must issue the operation $unlock(X)$ after all $read$ and $write$ operations on X are completed in T
4. A transaction T should not issue a $read_lock(X)$ / $write_lock(X)$ operation if it already holds a $read_lock$ or a $write_lock$ on item X
5. A transaction T should not issue an $unlock(X)$ operation unless it already holds a $read_lock$ or a $write_lock$ on item X

Conversion of locks

- Sometimes it is desirable to relax the rule 4 under *certain cenarios*
- That is, a transaction is already holding a lock on an item X and the transaction is allowed to ***convert*** the lock from one locked state to another
- Two types of lock conversions
 1. Upgarde a lock
 2. Downgrade a lock

Upgrading and downgrading of locks

- **Upgrading a lock:** A transaction T first issued a $read_lock(X)$ and later issues a $write_lock(X)$
 - This upgradation is possible *only if* T is the only transaction holding a read lock on X at the time it issues the $write_lock(X)$ operation; otherwise, T must wait
- **Downgrading a lock:** A transaction T first issued a $write_lock(X)$ and later issues a $read_lock(X)$
- When upgrading and downgrading of locks is used, the lock table must be updated accordingly

Do binary locks or read/write locks guarantee serializability?

T_1	T_2
<code>read_lock(Y);</code> <code>read_item(Y);</code> <code>unlock(Y);</code> <code>write_lock(X);</code> <code>read_item(X);</code> <code>X := X + Y;</code> <code>write_item(X);</code> <code>unlock(X);</code>	<code>read_lock(X);</code> <code>read_item(X);</code> <code>unlock(X);</code> <code>write_lock(Y);</code> <code>read_item(Y);</code> <code>Y := X + Y;</code> <code>write_item(Y);</code> <code>unlock(Y);</code>

Initial values: $X=20$, $Y=30$

Result serial schedule T_1
followed by T_2 : $X=50$, $Y=80$

Result of serial schedule T_2
followed by T_1 : $X=70$, $Y=50$

T_1	T_2
<code>read_lock(Y);</code> <code>read_item(Y);</code> <code>unlock(Y);</code> <code>write_lock(X);</code> <code>read_item(X);</code> <code>X := X + Y;</code> <code>write_item(X);</code> <code>unlock(X);</code>	<code>read_lock(X);</code> <code>read_item(X);</code> <code>unlock(X);</code> <code>write_lock(Y);</code> <code>read_item(Y);</code> <code>Y := X + Y;</code> <code>write_item(Y);</code> <code>unlock(Y);</code>

Result of schedule S:
 $X=50$, $Y=50$
(nonserializable)

The two-phase locking protocol

- To guarantee serializability, we must follow *an additional protocol* concerning the positioning of locking and unlocking operations in every transaction
- A transaction is said to follow the *two phase locking protocol* if all *read_lock* and *write_lock* operations precede the *first unlock* operation in the transaction

The two phases

- A transaction following the two-phase locking protocol has two phases
 1. **Growing phase:** During this phase, new locks on items can be acquired but none can be released
 2. **Shrinking phase:** During this phase, existing locks can be released but no new locks can be acquired
- If lock conversion is allowed, then upgrading of locks must be done during the expanding phase, and downgrading of locks must be done in the shrinking phase

Example

T_1	T_2
read_lock(Y); read_item(Y); unlock(Y); write_lock(X); read_item(X); $X := X + Y$; write_item(X); unlock(X);	read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); $Y := X + Y$; write_item(Y); unlock(Y);

T_1'
read_lock(Y); read_item(Y); write_lock(X); unlock(Y) read_item(X); $X := X + Y$; write_item(X); unlock(X);

T_2'
read_lock(X); read_item(X); write_lock(Y); unlock(X) read_item(Y); $Y := X + Y$; write_item(Y); unlock(Y);

T_1	T_2
read_lock(Y); read_item(Y); unlock(Y); write_lock(X); read_item(X); $X := X + Y$; write_item(X); unlock(X);	 read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); $Y := X + Y$; write_item(Y); unlock(Y);

The 2PL scheme guarantees serializability

- If *every* transaction in a schedule follows the two-phase locking protocol, the schedule is *guaranteed to be serializable*
- On contrary, assume 2PL doesn't ensure serializability
- Let T_1, T_2, \dots, T_n be the transactions which obey 2PL scheme and produce a non-serializable schedule
- Implies that, there must be a cycle in the precedence graph and let this cycle be (WLG) $T_n \rightarrow T_1 \rightarrow T_2 \rightarrow \dots \rightarrow T_{n-1} \rightarrow T_n$
- Let t_i be the time at which T_i obtains its *last* lock, for all $1 \leq i \leq n$
- For all the transactions $T_i \rightarrow T_j$, it must be that $t_i < t_j$
- As there is a cycle, we get $t_n < t_1 < t_2 < \dots < t_{n-1} < t_n$

Thank you!