

# Database Management Systems (DBMS)

Lec 27: Transaction Processing, Concurrency Control, and Recovery (Contd.)

Ramesh K. Jallu  
IIIT Raichur

Date: 18/06/21

# Recap

- Cascadeless schedules
- Concurrency control schemes
  - Lock based protocols
    - Binary locks
    - Shared/Exclusive locks (Read/ Write locks)
    - The two-phase locking protocol

# Today's plan

- Variants of the two-phase locking protocol
  - Conservative 2PL scheme
  - Strict 2PL scheme
  - Rigorous 2PL scheme
- Timestamp based protocols

# Conservative Two-Phase Locking

- A **conservative 2PL** (or **static 2PL**) requires a transaction to lock all the items it accesses *before the transaction begins execution*, by predeclaring its *read-set* and *write-set*
- If any of the predeclared items needed cannot be locked, the transaction does not lock any item; instead, *it waits until all the items are available for locking*
- Conservative 2PL is a *deadlock-free protocol*
- However, it is difficult to use in practice because of the need to predeclare the *read-set* and *write-set*, which may not be feasible

# Strict Two-Phase Locking

- A *strict schedule* is a schedule in which transactions can *neither read nor write* an item  $X$  until the last transaction that wrote  $X$  has committed (or aborted)
- A strict schedule is a cascadeless schedule
- In strict 2PL scheme, a transaction  $T$  does not release any of its exclusive (write) locks until *after* it commits or aborts
- Thus, no other transaction can read or write an item that is written by  $T$  unless  $T$  has committed
- Strict 2PL ensures strict schedule, but is *not a deadlock-free* protocol

# Rigorous Two-Phase Locking

- In rigorous 2PL scheme, a transaction  $T$  does not release any of its locks (exclusive or shared) until after it commits or aborts
- Like in strict 2PL, no other transaction can read or write an item that is written by  $T$  unless  $T$  has committed
- Rigorous 2PL is *not a deadlock-free* protocol
- It ensures a strict schedule

# Disadvantage of 2PL scheme

- WKT, locking combined with the 2PL protocol guarantees serializability of schedules
- The equivalent serial schedule is based on the order in which executing transactions lock the items they acquire
- If a transaction needs an item that is already locked, it may be forced to wait until the item is released
- Some transactions may be aborted and restarted because of the deadlock problem

# Concurrency Control Based on Timestamp Ordering

- Another approach to concurrency control involves using transaction timestamps to order transaction execution for an equivalent serial schedule
- Serializability is enforced by ordering conflicting operations in different transactions based on the transaction timestamps
- Concurrency control techniques based on timestamp ordering do not use locks; hence, *deadlocks cannot occur*



# Timestamps

- A *timestamp* is a unique identifier for each transaction  $T_i$  in the system, denoted by  $TS(T_i)$
- Typically, timestamp values are assigned in the order in which the transactions start execution
- If a transaction  $T_i$  has been assigned timestamp  $TS(T_i)$ , and a new transaction  $T_j$  enters the system, then  $TS(T_i) < TS(T_j)$
- Two methods to generate timestamps
  1. System clock
  2. Logical counter

# The Timestamp-Ordering Protocol (TOP)

- The TOP ensures that that any conflicting *read* and *write* operations are executed in timestamp order
- As a result, unlike 2PL, equivalent serial schedule obtained has the transactions in order of their timestamp values
- The protocol associates with each database item  $X$  two timestamp values
  - *read\_TS*( $X$ )
  - *write\_TS*( $X$ )

# Contd.

- When a transaction  $T_i$  issues a *read\_item(X)* or a *write\_item(X)* operation, the **TOP** compares the timestamp of  $T_i$  with *read\_TS(X)* and *write\_TS(X)* to ensure that the timestamp order of transaction execution is not violated
- If this order is violated, then transaction  $T_i$  is aborted and resubmitted to the system as a new transaction with a *new timestamp*
- This protocol *suffers* from cascading rollback as the schedules produced are not guaranteed to be recoverable
- However, by enforcing an additional protocol we can ensure that the schedules are cascadeless and recoverable

# Contd.

- When a transaction  $T_i$  issues the  $write\_item(X)$ 
  - If  $read\_TS(X) > TS(T_i)$  or if  $write\_TS(X) > TS(T_i)$ , then abort and roll back  $T_i$  and reject the operation
  - Otherwise, the system executes the  $write\_item(X)$  operation of  $T_i$  and sets  $write\_TS(X)$  to  $TS(T_i)$
- When a transaction  $T_i$  issues the  $read\_item(X)$ 
  - If  $write\_TS(X) > TS(T_i)$ , then abort and roll back  $T_i$  and reject the operation
  - Otherwise, the system executes the  $read\_item(X)$  operation of  $T_i$  and sets  $read\_TS(X)$  to the *larger* of  $TS(T_i)$  and the current  $read\_TS(X)$

# Example

$T_1$ : read( $B$ );  
read( $A$ );  
display( $A + B$ ).

$T_2$ : read( $B$ );  
 $B := B - 50$ ;  
write( $B$ );  
read( $A$ );  
 $A := A + 50$ ;  
write( $A$ );  
display( $A + B$ ).

$T_1$	$T_2$
read( $B$ )	read( $B$ ) $B := B - 50$ write( $B$ )
read( $A$ )	read( $A$ )
display( $A + B$ )	$A := A + 50$ write( $A$ ) display( $A + B$ )

# Observations

- Whenever the TO algorithm detects two *conflicting operations* that occur in the incorrect order, it rejects the later of the two operations by aborting the transaction that issued it
- The schedules produced by TO are hence guaranteed to be *conflict serializable* because conflicting operations are processed in timestamp order
- As mentioned earlier, deadlock does not occur with timestamp ordering
- However, cyclic restart (and hence starvation) may occur if a transaction is continually aborted and restarted

# Strict Timestamp Ordering Protocol (STOP)

- This variation of TOP, called STOP, ensures that the schedules are both *recoverable* and *serializable*
- Uses a limited form of locking, whereby reads of uncommitted items are postponed until the transaction that updated the item commits
- A transaction  $T$  issues a *read\_item(X)* or *write\_item(X)* such that  $TS(T) > write\_TS(X)$  has its read or write operation *delayed* until the transaction  $T'$  that *wrote* the value of  $X$  has committed or aborted
  - The item  $X$  is locked by transaction  $T'$  until  $T'$  is either committed or aborted
- The STOP *does not cause deadlock*, since  $T$  waits for  $T'$  only if  $TS(T) > TS(T')$

# Thomas's Write Rule

- It is a modified TOP, which allows greater potential concurrency than TOP
- When  $T_{27}$  attempts its  $write(Q)$  operation, we find that  $TS(T_{27}) < write\_TS(Q)$ , since  $write\_TS(Q) = TS(T_{28})$ . Thus, the  $write(Q)$  by  $T_{27}$  is rejected and transaction  $T_{27}$  must be rolled back
- Although the rollback of  $T_{27}$  is required by the timestamp-ordering protocol, it is unnecessary
- Since  $T_{28}$  has already written  $Q$ , the value that  $T_{27}$  is attempting to write is one that will never need to be read

$T_{27}$	$T_{28}$
read( $Q$ )	write( $Q$ )
write( $Q$ )	



# Thomas's Write Rule (Contd.)

- Any transaction  $T_i$  with  $TS(T_i) < TS(T_{28})$  that attempts a  $read(Q)$  will be rolled back, since  $TS(T_i) < write\_TS(Q)$
- Any transaction  $T_j$  with  $TS(T_j) > TS(T_{28})$  must read the value of  $Q$  written by  $T_{28}$ , rather than the value that  $T_{27}$  is attempting to write
- Thomas's write rule
  1. If  $read\_TS(X) > TS(T)$ , then abort and roll back  $T$  and reject the operation.
  2. If  $write\_TS(X) > TS(T)$ , then do not execute the write operation but continue processing
  3. If neither the condition in 1 nor the condition in 2 occurs, then execute the  $write\_item(X)$  operation of  $T$  and set  $write\_TS(X)$  to  $TS(T)$

Thank you!